

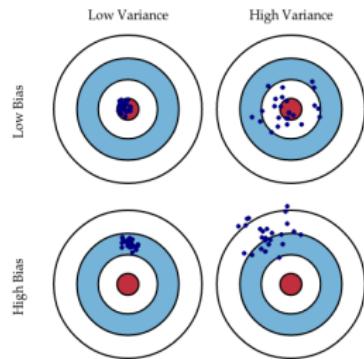
Reinforcement Learning

RL in finite MDPs



Marcello Restelli

February, 2024





RL techniques

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Model-free vs Model-based
- On-policy vs Off-policy
- Online vs Offline
- Tabular vs Function Approximation
- Value-based vs Policy-based vs Actor-Critic



RL problems

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- **Model-free Prediction:** Estimate the value function of an **unknown** MRP (MDP + policy)
- **Model-free Control:** Optimize the value function of an **unknown** MDP



Monte–Carlo Reinforcement Learning

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- MC methods learn **directly** from episodes of **experience**
- MC is **model-free**: no knowledge of MDP transitions/rewards
- MC learns from **complete** episodes: no bootstrapping
- MC uses the simplest possible idea: **value = mean return**
- Caveat: can only apply MC to **episodic** MDPs
 - All episodes **must terminate**



Monte Carlo for Prediction and Control

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- MC can be used for **prediction**:
 - **Input:** Episodes of experience $\{s_1, a_1, r_2, \dots, s_T\}$ generated by following policy π in given MDP
 - or: Episodes of experience $\{s_1, a_1, r_2, \dots, s_T\}$ generated by MRP
 - **Output:** Value function V^π
- Or for **control**:
 - **Input:** Episodes of experience $\{s_1, a_1, r_2, \dots, s_T\}$ in given MDP
 - **Output:** Optimal value function V^*
 - **Output:** Optimal policy π^*



Estimation of Mean: Monte Carlo

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Let X be a random variable with mean $\mu = \mathbb{E}[x]$ and variance $\sigma^2 = \text{Var}[X]$. Let $x_i \sim X$, $i = 1, \dots, n$ be n i.i.d. realizations of X .
- Empirical mean** of X :

$$\hat{\mu}_n = \frac{1}{n} \sum_{i=1}^n x_i$$

- We have $\mathbb{E}[\hat{\mu}_n] = \mu$, $\text{Var}[\hat{\mu}_n] = \frac{\text{Var}[x]}{n}$
 - Weak law of large numbers:** $\hat{\mu}_n \xrightarrow{P} \mu$
 $(\lim_{n \rightarrow \infty} \mathbb{P}(|\hat{\mu}_n - \mu| > \epsilon) = 0)$
 - Strong law of large numbers:** $\hat{\mu}_n \xrightarrow{a.s.} \mu$
 $(\mathbb{P}(\lim_{n \rightarrow \infty} \hat{\mu}_n = \mu) = 1)$
 - Central limit theorem:** $\sqrt{n}(\hat{\mu}_n - \mu) \xrightarrow{D} \mathcal{N}(0, \text{Var}[x])$



Monte–Carlo Policy Evaluation

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- **Goal:** learn V^π from experience under policy π

$$s_1, a_1, r_2, \dots, s_T \sim \pi$$

- Recall that the **return** is the total discounted reward:

$$v_t = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{T-1} r_{t+T}$$

- Recall that the **value function** is the expected return:

$$V^\pi(s) = \mathbb{E}[v_t | s_t = s]$$

- Monte Carlo policy evaluation uses **empirical mean** return instead of expected return
 - **first visit:** average returns only for the first time s is visited (**unbiased** estimator)
 - **every visit:** average returns for every time s is visited (**biased** but **consistent** estimator)



First–Visit Monte–Carlo Policy Evaluation

Marcello
Restelli

Model–free
Prediction

Monte–Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model–free
Control

On–Policy
Monte–Carlo Control
On–Policy
Temporal–Difference
Learning

Off–Policy Learning

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state–value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

loop

Generate an episode using π

for each state s in the episode **do**

$R \leftarrow$ return following the first occurrence of s

Append R to $Returns(s)$

$V(s) \leftarrow$ average($Returns(s)$)

end for

end loop



Every–Visit Monte–Carlo Policy Evaluation

Marcello
Restelli

Model–free
Prediction

Monte–Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model–free
Control

On–Policy
Monte–Carlo Control

On–Policy
Temporal–Difference
Learning

Off–Policy Learning

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state–value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

loop

Generate an episode using π

for each state s in the episode **do**

for each occurrence of state s in the episode **do**

$R \leftarrow$ return following this occurrence of s

Append R to $Returns(s)$

$V(s) \leftarrow$ average($Returns(s)$)

end for

end for

end loop



First–Visit vs Every–Visit

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

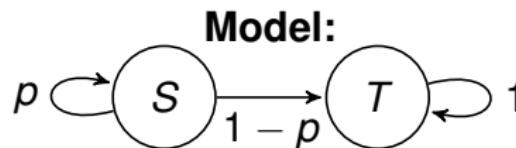
TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

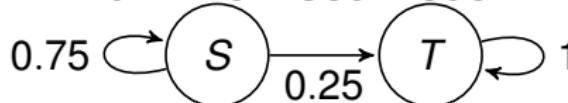


Reward is +1 on every step

Sample Path: $S \rightarrow S \rightarrow S \rightarrow S \rightarrow T$

$$V^{FV}(S) = 4 \quad V^{EV}(S) = 2.5$$

Max Likelihood Model:





First–Visit vs Every–Visit Crossover

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

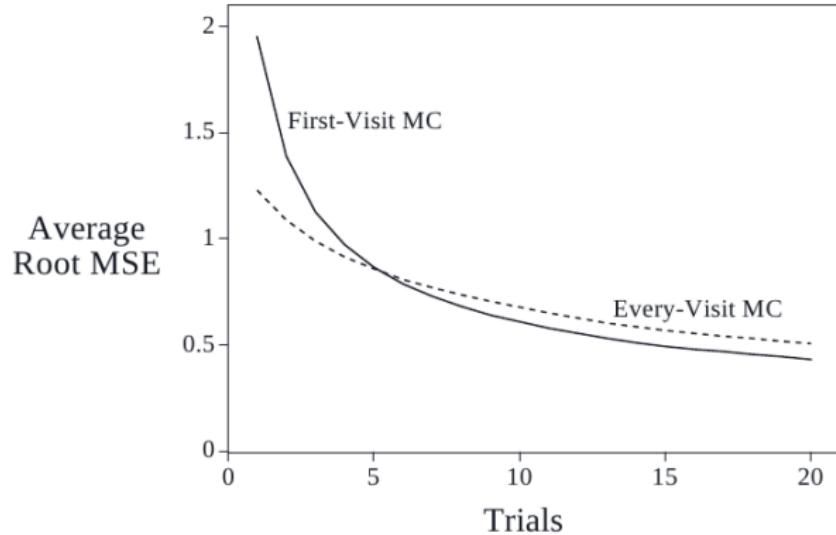
Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning





Blackjack Example

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- **Goal:** Have your card sum be greater than the dealers without exceeding 21
- **States** (200 of them):
 - current sum (12–21)
 - dealer's showing card (ace–10)
 - do I have a usable ace?
- **Reward:** +1 for winning, 0 for a draw, -1 for losing
- **Actions:** stand (stop receiving cards), hit (receive another card)
- **Policy:** Stand if my sum is 20 or 21, else hit



Blackjack Example

After Monte–Carlo Learning

Marcello
Restelli

Model-free
Prediction

Monte–Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

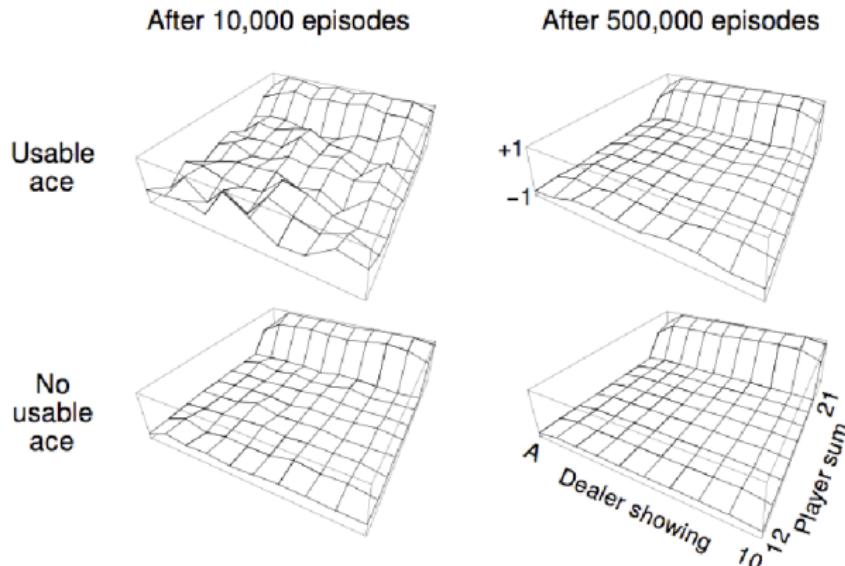
$TD(\lambda)$

Model-free
Control

On-Policy
Monte–Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning





Incremental Mean

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

The mean $\hat{\mu}_1, \hat{\mu}_2, \dots$ of a sequence x_1, x_2, \dots can be computed incrementally

$$\begin{aligned}\hat{\mu}_k &= \frac{1}{k} \sum_{j=1}^k x_j \\ &= \frac{1}{k} \left(x_k + \sum_{j=1}^{k-1} x_j \right) \\ &= \frac{1}{k} (x_k + (k-1)\hat{\mu}_{k-1}) \\ &= \hat{\mu}_{k-1} + \frac{1}{k} (x_k - \hat{\mu}_{k-1})\end{aligned}$$



Incremental Monte–Carlo Updates

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- Update $V(s)$ **incrementally** after episode $s_1, a_1, r_2, \dots, s_T$
- For each state s_t with return v_t

$$N(s_t) \leftarrow N(s_t) + 1$$

$$V(s_t) \leftarrow V(s_t) + \frac{1}{N(s_t)}(v_t - V(s_t))$$

- In **non-stationary** problems, it is useful to track a running mean, i.e., **forget** old episodes

$$V(s_t) \leftarrow V(s_t) + \alpha(v_t - V(s_t))$$



Stochastic Approximation

Estimation of Mean

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

Let X be a random variable in $[0, 1]$ with mean $\mu = \mathbb{E}[X]$. Let $x_i \sim X, i = 1, \dots, n$ be n i.i.d. realizations of X .

Consider the estimator (**exponential average**)

$$\mu_i = (1 - \alpha_i)\mu_{i-1} + \alpha_i x_i,$$

with $\mu_1 = x_1$ and α_i 's are **step-size parameters** or **learning rates**

Proposition

If $\sum_{i \geq 0} \alpha_i = \infty$ and $\sum_{i \geq 0} \alpha_i^2 < \infty$, then $\hat{\mu}_n \xrightarrow{\text{a.s.}} \mu$, i.e., the estimator $\hat{\mu}_n$ is **consistent**

Note: The step sizes $\alpha_i = \frac{1}{i}$ satisfy the above conditions. In this case, the exponential average gives us the empirical mean $\hat{\mu}_n = \frac{1}{n} \sum_{i=1}^n x_i$, which is **consistent** according to the **strong law of large numbers**



Monte–Carlo Backups

Marcello
Restelli

Model-free
Prediction

Monte–Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte–Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- **Entire episode** included
- Only **one choice** at each state (unlike DP)
- MC **does not bootstrap**
- Time required to estimate one state **does not depend** on the total number of states



Temporal Difference Learning

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- TD methods **learn directly** from episodes of experience
- TD is **model-free**: no knowledge of MDP transitions/rewards
- TD learns from **incomplete** episodes: **bootstrapping**
- TD updates a **guess** towards a **guess**



TD Prediction

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- **Goal:** learn V^π online from experience under policy π
- **Recall:** incremental every-visit Monte Carlo

$$V(s_t) \leftarrow V(s_t) + \alpha(r_t - V(s_t))$$

- **Simplest** temporal-difference learning algorithm:
TD(0)

- Update value $V(s_t)$ towards **estimated return**

$$r_{t+1} + \gamma V(s_{t+1})$$

$$V(s_t) \leftarrow V(s_t) + \alpha(r_{t+1} + \gamma V(s_{t+1}) - V(s_t))$$

- $r_{t+1} + \gamma V(s_{t+1})$ is called the **TD target**
- $\delta_t = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$ is called the **TD error**



Bias–Variance Trade–Off

Conceptual Definition

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- **Error due to bias:** is the difference between the expected prediction of our model and the actual value we want to predict
- **Error due to variance:** is the variability of a model prediction for a given data point



Bias–Variance Trade–Off

Graphical Definition

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

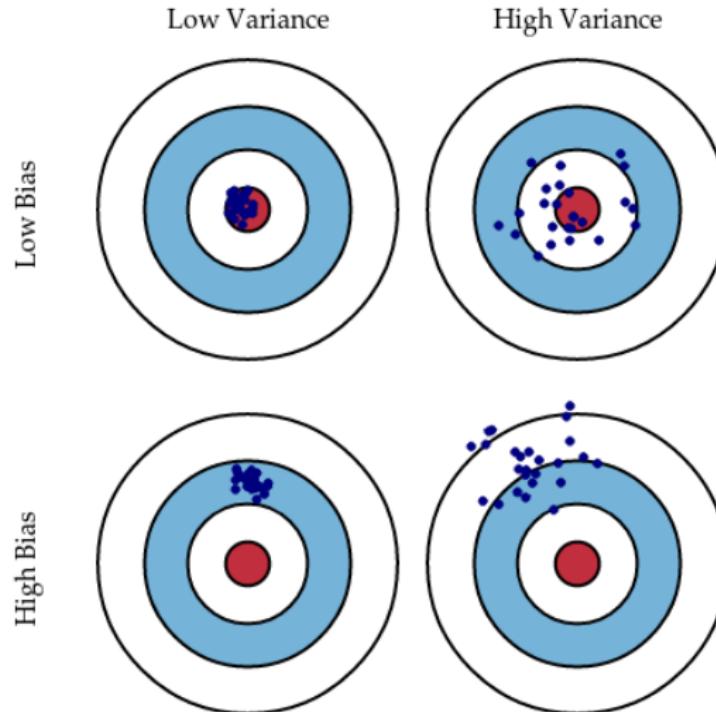
Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning



Bias–Variance Trade–Off

Mathematical Definition

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

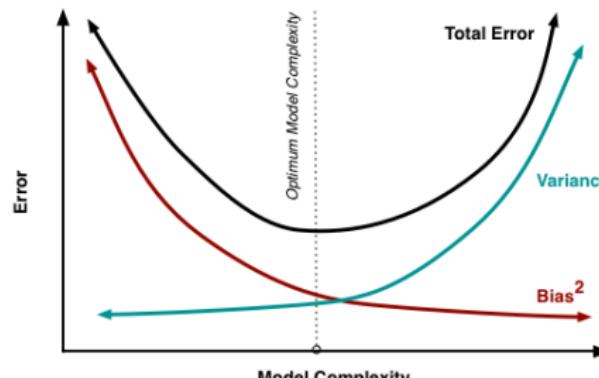
- We want to **predict** $Y = f(X) + \epsilon$, where the error term ϵ is normally distributed $\epsilon \sim \mathcal{N}(0, \sigma_\epsilon)$
- We estimate a model $\hat{f}(X)$, with the following **prediction error** at point x :

$$Err(x) = \mathbb{E}[(Y - \hat{f}(x))^2]$$

- This error can be decomposed into **bias** and **variance**:

$$Err(x) = \left(f(x) - \mathbb{E}[\hat{f}(x)] \right)^2 + \mathbb{E} \left[\mathbb{E}[\hat{f}(x)] - \hat{f}(x) \right]^2 + \sigma_\epsilon^2$$

$$Err(x) = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$





Bias–Variance Trade–Off

MC vs TD

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $\text{TD}(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Return $v_t = r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^{T-1} r_{t+T}$ is an **unbiased** estimate of $V^\pi(s_t)$
- TD target $r_{t+1} + \gamma V(s_{t+1})$ is a **biased** estimate of $V^\pi(s_t)$
 - Unless $V(s_{t+1}) = V^\pi(s_{t+1})$
- But the TD target is much **lower variance**:
 - Return depends on **many** random actions, transitions, rewards
 - TD target depends on **one** random action, transition, reward



Bias–Variance comparison between MC and TD

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- MC has high variance, zero bias
 - **Good convergence** properties
 - Works well with **function approximation**
 - **Not** very sensitive to **initial value**
 - Very **simple** to understand and use
- TD has low variance, some bias
 - Usually **more efficient** than MC
 - TD(0) converges to $V^\pi(s)$
 - **Problem** with function approximation
 - More **sensitive** to initial values

Random Walk Example

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

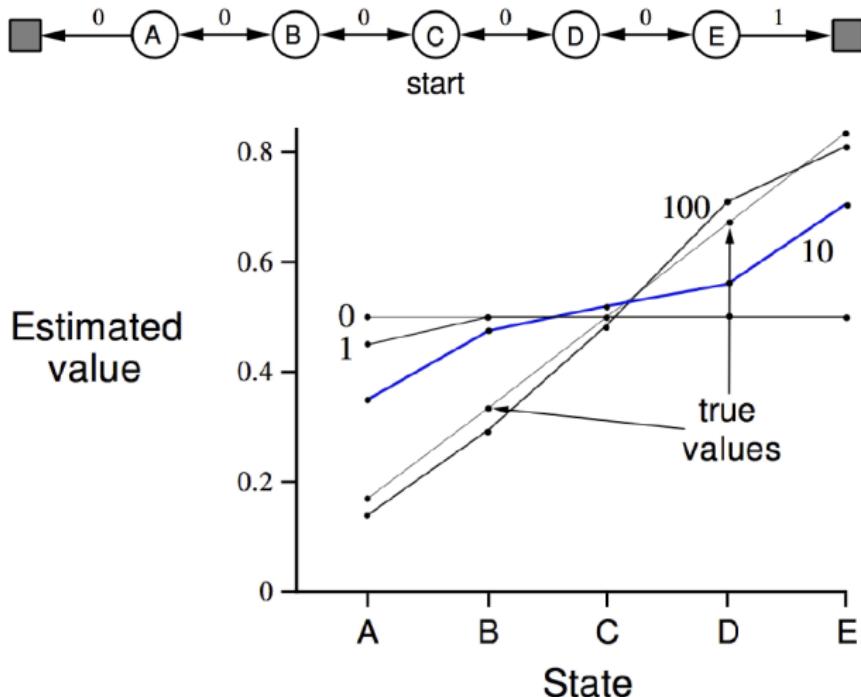
TD(λ)

Model-free Control

On-Policy Monte-Carlo Control

On-Policy Temporal-Difference Learning

Off-Policy Learning





Random Walk

MC vs TD

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

TD(λ)

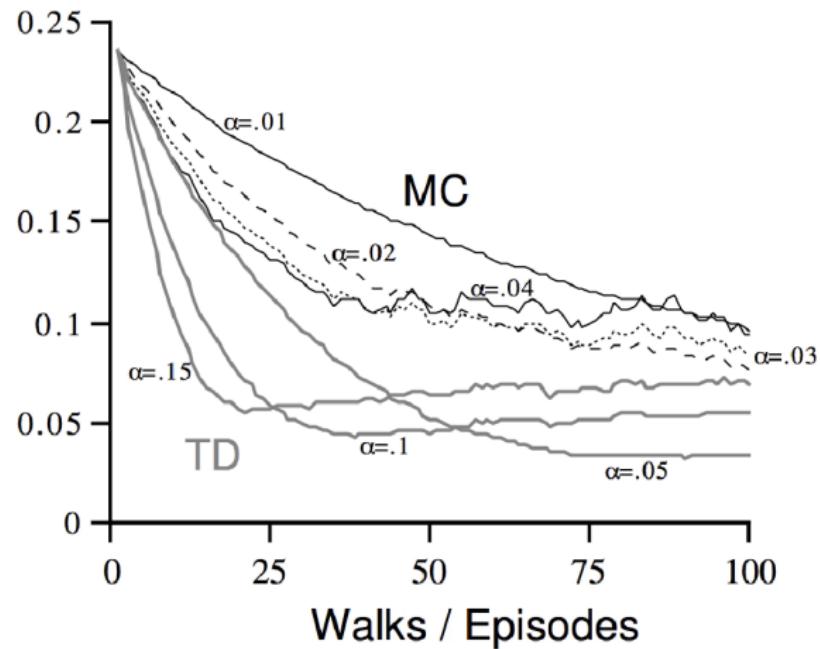
Model-free Control

On-Policy Monte-Carlo Control

On-Policy Temporal-Difference Learning

Off-Policy Learning

RMS error,
averaged
over states





Batch MC and TD

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- MC and TD **converge**: $V(s) \rightarrow V^\pi(s)$ as experience $\rightarrow \infty$
- But what about **batch** solution for finite experience?

$$s_1^1, a_1^1, r_2^1, \dots, s_{T_1}^1$$

⋮

$$s_1^K, a_1^K, r_2^K, \dots, s_{T_1}^K$$

- e.g., **repeatedly** sample episode $k \in [1, K]$
- Apply MC or TD(0) to episode k



AB Example

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- Two states A, B
- Undiscounted
- 8 episodes of experience

$A, 0, B, 0$

$B, 1$

$B, 1$

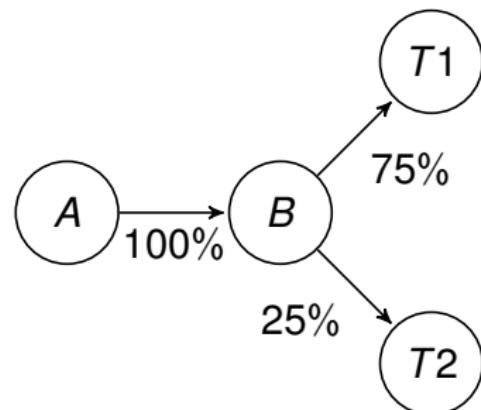
$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 0$



- What is $V(A), V(B)$?



Certainty Equivalence

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- MC converges to solution with **minimum mean-squared error**
 - Best fit to the **observed returns**

$$\sum_{k=1}^K \sum_{t=1}^T (v_t^k - V(s_t^k))^2$$

- In the AB example, $V(A) = 0$
- TD(0) converges to solution of **max likelihood Markov model**
 - Solution to the MDP $\langle \mathcal{S}, \mathcal{A}, \hat{P}, \hat{R}, \gamma, \mu \rangle$ that **best fits the data**

$$\hat{P}(s'|s, a) = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{t=1}^T \mathbf{1}(s_t^k, a_t^k, s_{t+1}^k = s, a, s')$$

$$\hat{R}(s, a) = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{t=1}^T \mathbf{1}(s_t^k, a_t^k = s, a) r_t^k$$

- In the AB example, $V(A) = 0.75$



Comparison between MC and TD

Markov Property

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- TD **exploits** Markov property
 - Usually more efficient in **Markov environments**
- MC **does not exploit** Markov property
 - Usually more efficient in **non-Markov environments**



MC vs TD vs DP

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

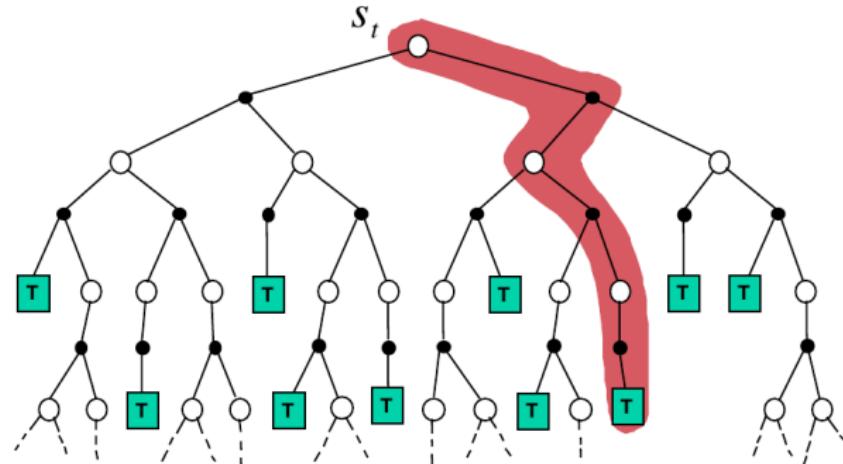
On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

$$V(s_t) \leftarrow V(s_t) + \alpha[v_t - V(s_t)]$$

where R_t is the actual return following state s_t





MC vs TD vs DP

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

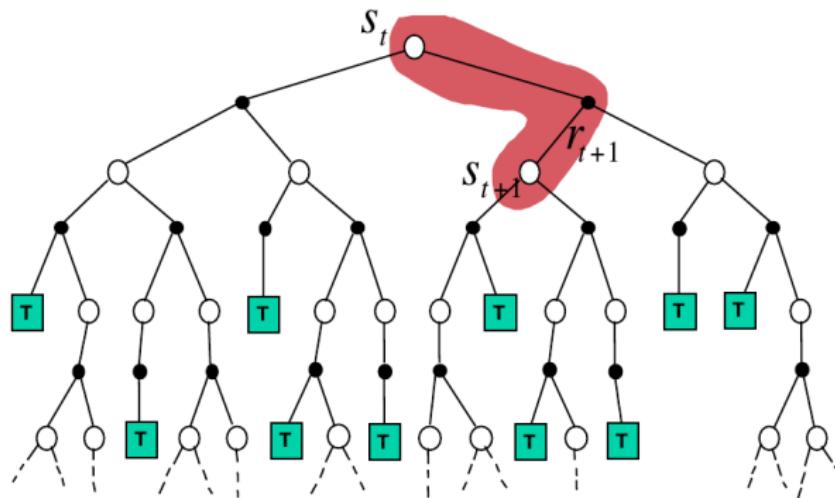
TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

$$V(s_t) \leftarrow V(s_t) + \alpha[r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

where R_t is the actual return following state s_t





MC vs TD vs DP

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

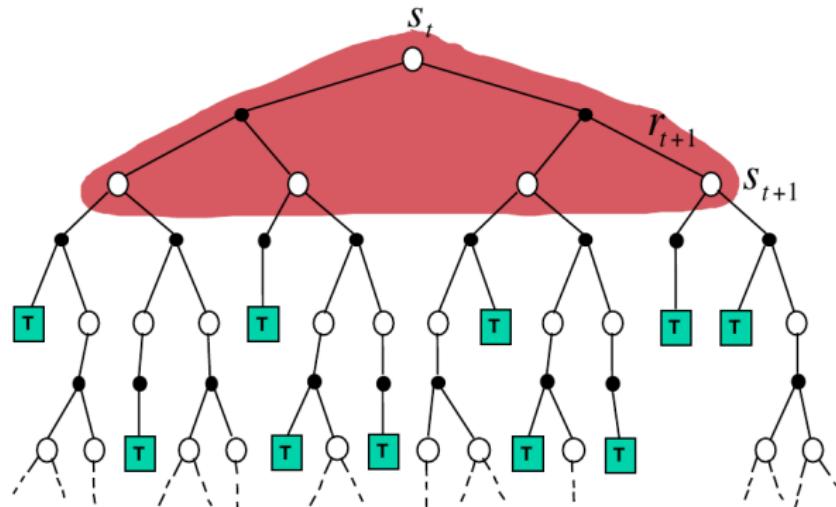
TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

$$V(s_t) \leftarrow \mathbb{E}_\pi[r_{t+1} + \gamma V(s_{t+1})]$$

where R_t is the actual return following state s_t





MC vs TD vs DP

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

	Bootstrapping	Sampling
MC	No	Yes
TD	Yes	Yes
DP	Yes	No



MC vs TD vs DP

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

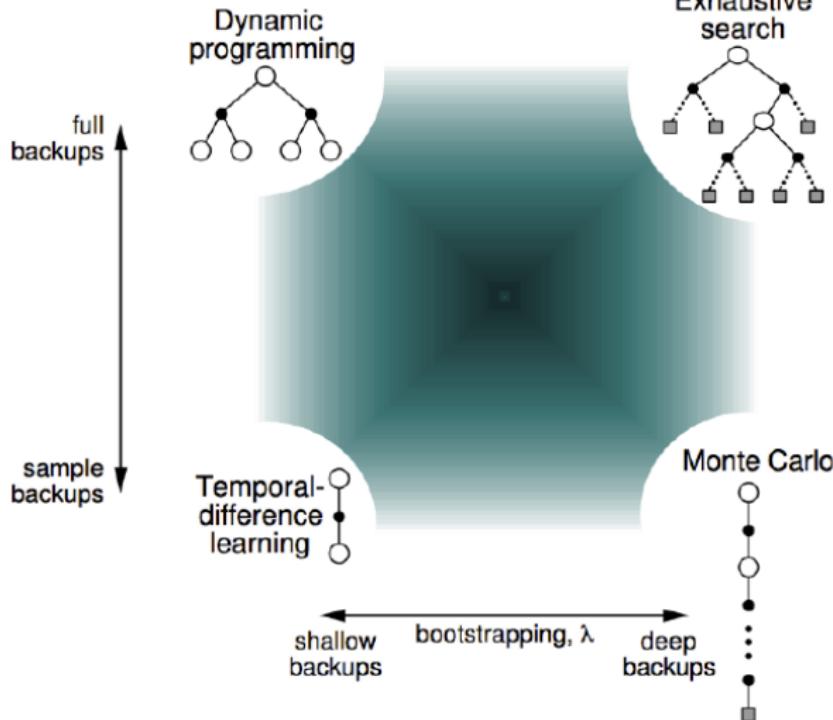
Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning





n -Step Prediction

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

$\text{TD}(\lambda)$

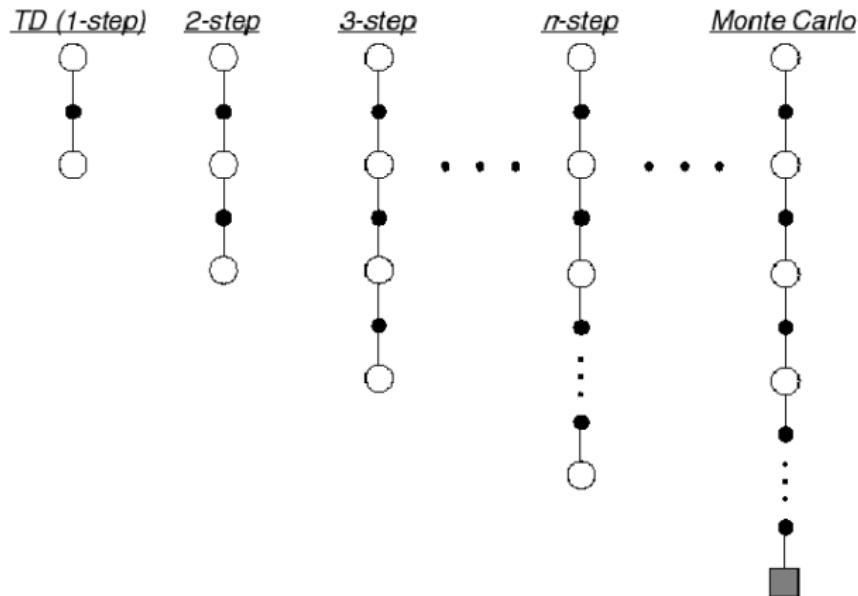
Model-free Control

On-Policy Monte-Carlo Control

On-Policy Temporal-Difference Learning

Off-Policy Learning

Let TD target look n steps into the future





n -Step Return

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $\text{TD}(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Consider the following n -step returns for
 $n = 1, 2, \dots, \infty$:

$$n = 1 \quad (\text{TD}) \quad v_t^{(1)} = r_{t+1} + \gamma V(s_{t+1})$$

$$n = 2 \quad v_t^{(2)} = r_{t+1} + \gamma r_{t+2} + \gamma^2 V(s_{t+2})$$

$$\vdots \qquad \vdots$$

$$n = \infty \quad (\text{MC}) \quad v_t^{(\infty)} = r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^{T-1} r_T$$

- Define the n -step return

$$v_t^{(n)} = r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^{n-1} r_{t+n} + \gamma^n V(s_{t+n})$$

- n -step temporal-difference learning

$$V(s_t) \leftarrow V(s_t) + \alpha(v_t^{(n)} - V(s_t))$$



Large Random Walk Example

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

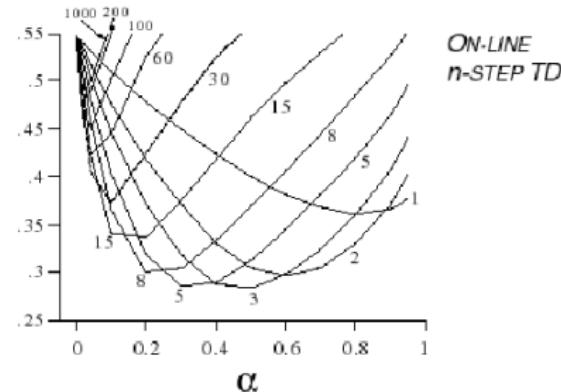
Comparison between MC and TD

$\text{TD}(\lambda)$

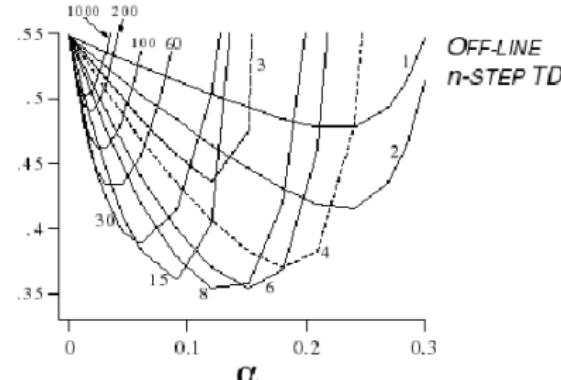
Model-free Control

On-Policy Monte-Carlo Control
On-Policy Temporal-Difference Learning
Off-Policy Learning

RMS error,
averaged over
first 10 episodes



RMS error,
averaged over
first 10 episodes





Averaging n -step Returns

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- We can **average** n -step returns over **different n**
- e.g., average the 2-step and 4-step returns

$$\frac{1}{2}v^{(2)} + \frac{1}{2}v^{(4)}$$

- **Combines information** from two different time-steps
- Can we **efficiently** combine information from all time-steps?



λ -return

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $\text{TD}(\lambda)$

Model-free
Control

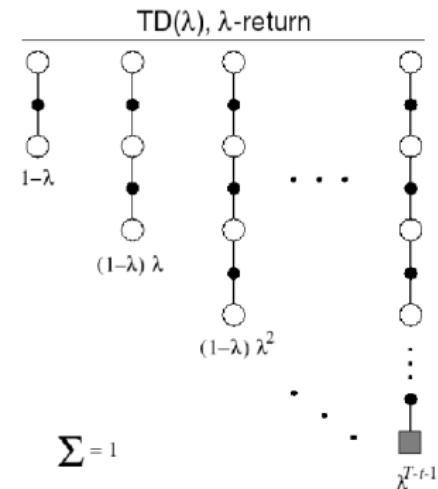
On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- The λ -return v_t^λ combines all n -step returns $v_t^{(n)}$
- Using **weight** $(1 - \lambda)\lambda^{n-1}$

$$v_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} v_t^{(n)}$$

- **Forward-view TD(λ)**

$$V(s_t) \leftarrow V(s_t) + \alpha \left(v_t^\lambda - V(s_t) \right) \quad \Sigma = 1$$





TD(λ) Weighting Function

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

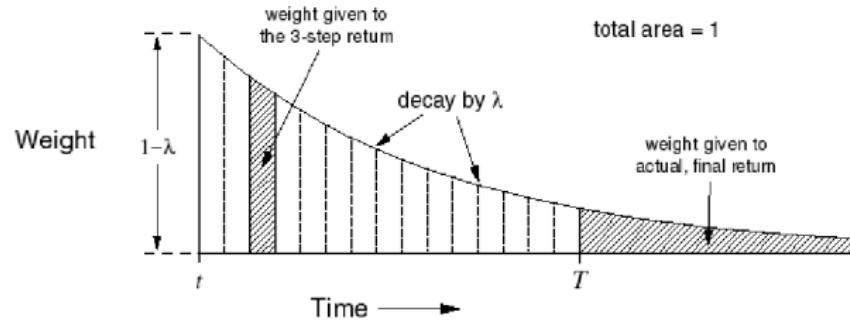
Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning



$$v_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} v_t^{(n)}$$

Forward-view TD(λ)

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

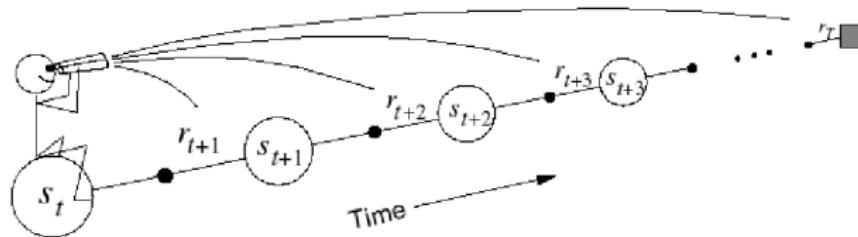
Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning



- Update value function towards the λ -return
- Forward-view looks into the **future** to compute v_t^λ
- Like MC, can only be computed from **complete episodes**



Forward–view TD(λ) on Large Random Walk

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

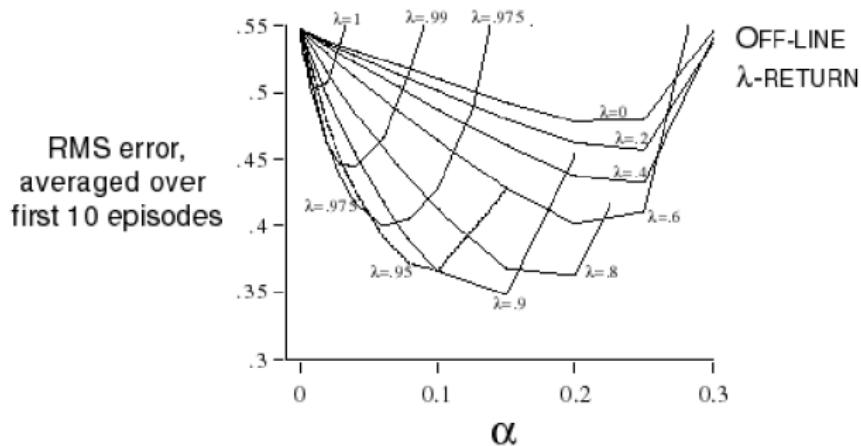
Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning





Backward–view TD(λ)

Marcello
Restelli

Model–free
Prediction

Monte–Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model–free
Control

On–Policy
Monte–Carlo Control

On–Policy
Temporal–Difference
Learning

Off–Policy Learning

- Forward view provides **theory**
- Backward view provides **mechanism**
- Update **online**, every step, from **incomplete sequences**

Eligibility Traces

Marcello Restelli

Model-free Prediction

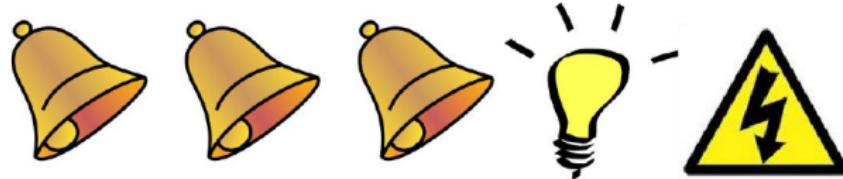
Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning
Comparison between MC and TD

TD(λ)

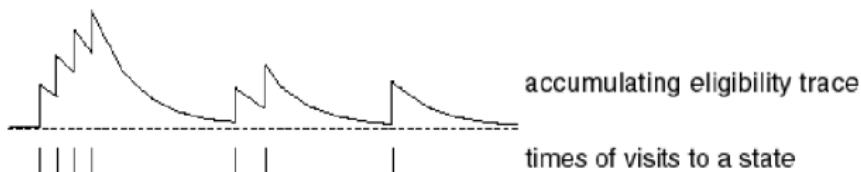
Model-free Control

On-Policy Monte-Carlo Control
On-Policy Temporal-Difference Learning
Off-Policy Learning



- **Credit assignment problem:** did bell or light cause shock?
- **Frequency heuristic:** assign credit to the most frequent states
- **Recency heuristics:** assign credit to the most recent states
- **Eligibility traces** combine both heuristics

$$e_{t+1}(s) = \gamma \lambda e_t(s) + \mathbf{1}(s = s_t)$$



Backward-view TD(λ)

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

TD(λ)

Model-free Control

On-Policy Monte-Carlo Control

On-Policy Temporal-Difference Learning

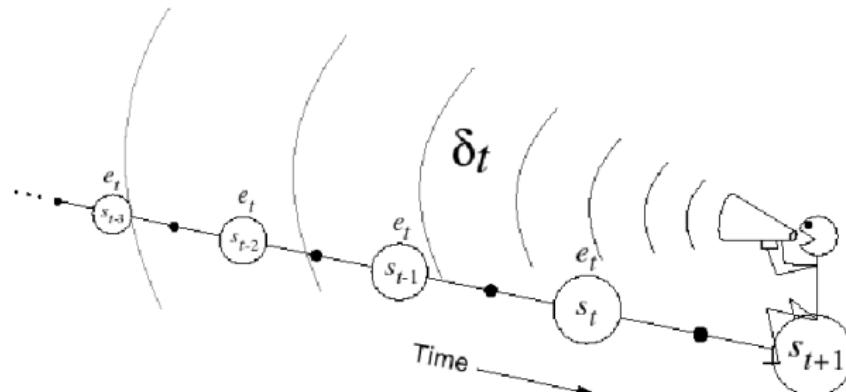
Off-Policy Learning

- **Update value $V(s)$ for every state s**
- In proportion to TD-error δ_t and **eligibility trace $e_t(s)$**

$$e_0(s) = 0$$

$$e_t(s) = \gamma \lambda e_{t-1}(s) + \mathbf{1}(s = s_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t e_t(s)$$





Backward-view TD(λ) Algorithm

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Initialize $V(s)$ arbitrarily

for all episodes **do**

$$e(s) = 0, \quad \forall s \in \mathcal{S}$$

Initialize s

repeat

$a \leftarrow$ action given by π for s

Take action a , observe reward r , and next state s'

$$\delta \leftarrow r + \gamma V(s') - V(s)$$

$$e(s) \leftarrow e(s) + 1$$

for all $s \in \mathcal{S}$ **do**

$$V(s) \leftarrow V(s) + \alpha \delta e(s)$$

$$e(s) \leftarrow \gamma \lambda e(s)$$

end for

$$s \leftarrow s'$$

until s is terminal

end for



TD(λ) and TD(0)

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- When $\lambda = 0$, only current state is updated

$$e_t(s) = \mathbf{1}(s = s_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t e_t(s)$$

- This is exactly equivalent to TD(0) update

$$V(s_t) \leftarrow V(s_t) + \alpha \delta_t$$



Telescoping in TD(1)

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

When $\lambda = 1$, sum of TD errors telescopes into MC error

$$\begin{aligned} & \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \cdots + \gamma^{T-t}\delta_{T-1} \\ = & r_{t+1} + \gamma V(s_{t+1}) - V(s_t) \\ + & \gamma r_{t+2} + \gamma^2 V(s_{t+2}) - \gamma V(s_{t+1}) \\ + & \gamma^2 r_{t+3} + \gamma^3 V(s_{t+3}) - \gamma^2 V(s_{t+2}) \\ & \vdots \\ + & \gamma^{T-1} r_{t+T} + \gamma^T V(s_{t+T}) - \gamma^{T-1} V(s_{t+T-1}) \\ = & r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots + \gamma^{T-1} r_{t+T} - V(s_t) \\ = & v_t - V(s_t) \end{aligned}$$



TD(λ) and TD(1)

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- TD(1) is **roughly equivalent** to every-visit Monte-Carlo
- Error is accumulated online, **step-by-step**
- If value function is only **updated offline** at end of episode, then the total update is **exactly** the same as MC
- If value function is **updated online** after every step, then TD(1) may have **different** total update to MC



Forwards and Backwards TD(λ)

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Consider an episode where s is **visited once** at time-step k
- TD(λ) eligibility trace **discounts time** since visit

$$e_t(s) = \gamma \lambda e_{t-1}(s) + \mathbf{1}(s_t = s)$$

$$= \begin{cases} 0 & \text{if } t < k \\ (\gamma \lambda)^{t-k} & \text{if } t \geq k \end{cases}$$

- Backward TD(λ) updates **accumulate error online**

$$\sum_{t=1}^T \alpha \delta_t e_t(s) = \alpha \sum_{t=k}^T (\gamma \lambda)^{t-k} \delta_t = \alpha(v_k - V(s_k))$$

- By end of episode it accumulates **total error** for λ -return
- For **multiple** visits to s , $e_t(s)$ accumulates **many errors**



Equivalence of Forward and Backward TD

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Theorem

The sum of offline updates is identical for forward-view and backward-view TD(λ)

$$\sum_{t=1}^T \alpha \delta_t e_t(s) = \sum_{t=1}^T \alpha(v_t^\lambda - V(s_t)) \mathbf{1}(s_t = s)$$

- In **practice**, value function is updated **online** by TD(λ)
- But if α is **small** then equivalence is almost exact

Replacing Traces

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

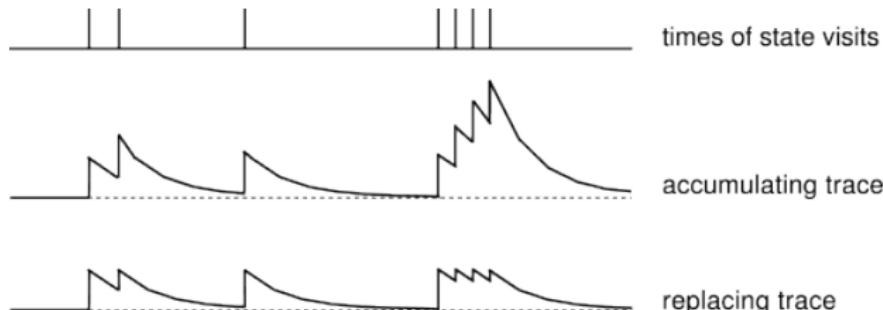
Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- Using **accumulating** traces, **frequently** visited states can have eligibilities greater than 1
 - This can be a **problem for convergence**
- Replacing traces:** Instead of adding 1 when you visit a state, set that trace to 1

$$e_t(s) = \begin{cases} \gamma \lambda e_{t-1}(s) & \text{if } s \neq s_t \\ 1 & \text{if } s = s_t \end{cases}$$





Use of Model-Free Control

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Some example problems that can be modeled as MDPs:

- Elevator
- Parallel Parking
- Ship Steering
- Bioreactor
- Helicopter
- Airplane Logistics
- Robocup Soccer
- Quake
- Portfolio management
- Protein folding
- Robot walking
- Game of Go

For most of these problems, either:

- MDP model is **unknown**, but experience can be **sampled**
- MDP model is **known**, but is **too big** to use, except by samples

Model-free control can solve these problems



On and Off-Policy Learning

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- **On-policy learning**

- “Learn on the job”
- Learn about policy π from experience sampled from π

- **Off-policy learning**

- “Learn over someone’s shoulder”
- Learn about policy π from experience sampled from $\bar{\pi}$



Generalized Policy Iteration (Refresher)

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

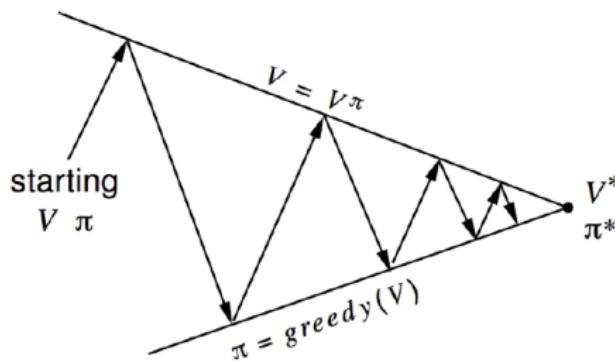
Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

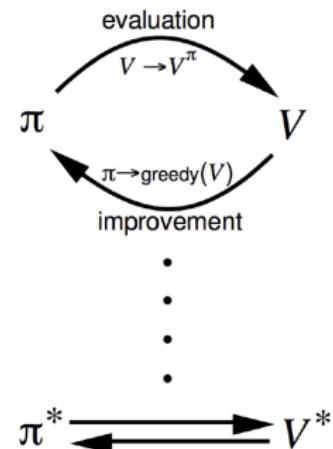
On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning



- **Policy evaluation:** Estimate V^{π}
 - e.g., Iterative policy evaluation
- **Policy improvement:** Generate $\pi' \geq \pi$
 - e.g., Greedy policy improvement





Generalized Policy Iteration with Monte–Carlo Evaluation

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

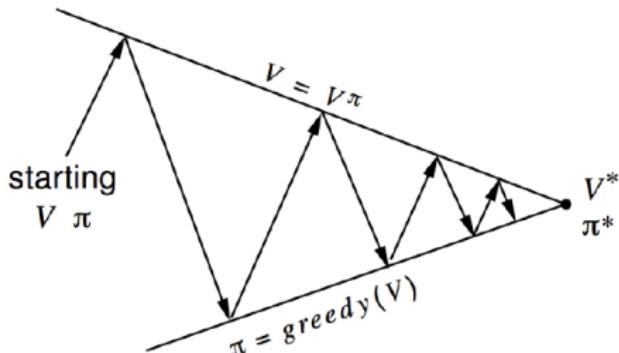
TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning



- **Policy Evaluation:** Monte–Carlo policy evaluation,
 $V = V^\pi?$
- **Policy Improvement:** Greedy policy improvement?



Model-Free Policy Iteration Using Action-Value Function

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning
Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control
On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Greedy policy improvement over $V(s)$ **requires model** of MDP

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} \left\{ R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s') \right\}$$

- Greedy policy improvement over $Q(s, a)$ is **model-free**

$$\pi'(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$$



Generalized Policy Iteration with Monte–Carlo Evaluation

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

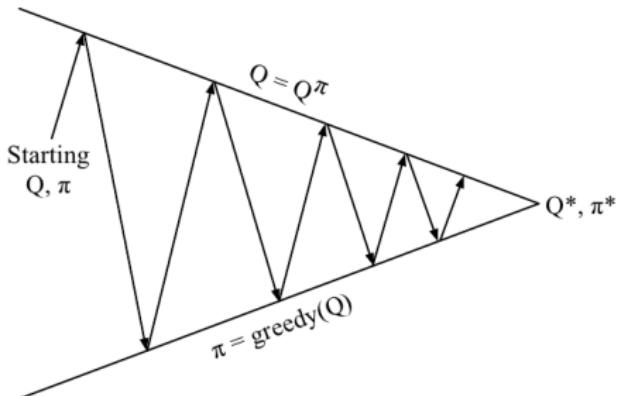
TD(λ)

Model-free Control

On-Policy Monte-Carlo Control

On-Policy Temporal-Difference Learning

Off-Policy Learning



- **Policy Evaluation:** Monte–Carlo policy evaluation,
 $Q = Q^\pi$
- **Policy Improvement:** Greedy policy improvement?



On-Policy Exploration

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning



"Behind one door is tenure - behind the other
is flipping burgers at McDonald's."

- There are two doors in front of you
- You open the left door and get reward 0, $V(left) = 0$
- You open the right door and get reward +1, $V(right) = +1$
- You open the right door and get reward +3, $V(right) = +2$
- You open the right door and get reward +2, $V(right) = +2$
- ⋮
- Are you sure you've chosen the **best** door?



ϵ -Greedy Exploration

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control
On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- **Simplest** idea for ensuring **continual exploration**
- **All** m actions are tried with **non-zero** probability
- With probability $1 - \epsilon$ choose the **greedy action**
- With probability ϵ choose an action **at random**

$$\pi(s, a) = \begin{cases} \frac{\epsilon}{m} + 1 - \epsilon & \text{if } a^* = \arg \max_{a \in \mathcal{A}} Q(s, a) \\ \frac{\epsilon}{m} & \text{otherwise} \end{cases}$$



ϵ -Greedy Policy Improvement

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Theorem

For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to Q^π is an improvement

$$\begin{aligned} Q^\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a|s) Q^\pi(s, a) \\ &= \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} Q^\pi(s, a) + (1 - \epsilon) \max_{a \in \mathcal{A}} Q^\pi(s, a) \\ &\geq \frac{\epsilon}{m} \sum_{a \in \mathcal{A}} Q^\pi(s, a) + (1 - \epsilon) \sum_{a \in \mathcal{A}} \frac{\pi(a|s) - \frac{\epsilon}{m}}{1 - \epsilon} Q^\pi(s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) Q^\pi(s, a) = V^\pi(s) \end{aligned}$$

Therefore from policy improvement theorem, $V^{\pi'(s)} \geq V^\pi(s)$



Monte–Carlo Policy Iteration

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

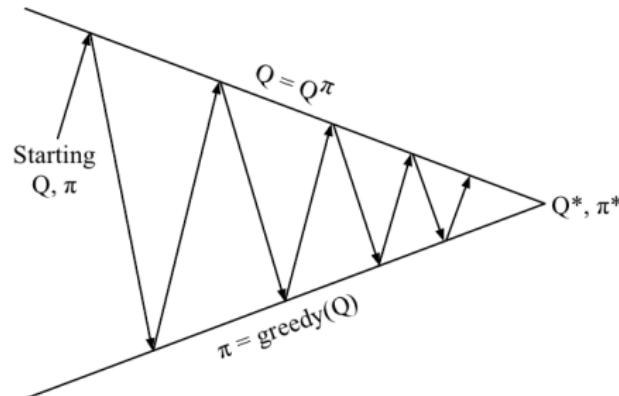
TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning



- **Policy Evaluation:** Monte–Carlo policy evaluation,
$$Q = Q^\pi$$
- **Policy Improvement:** ϵ –greedy policy improvement



Monte–Carlo Control

Marcello
Restelli

Model-free
Prediction

Monte–Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

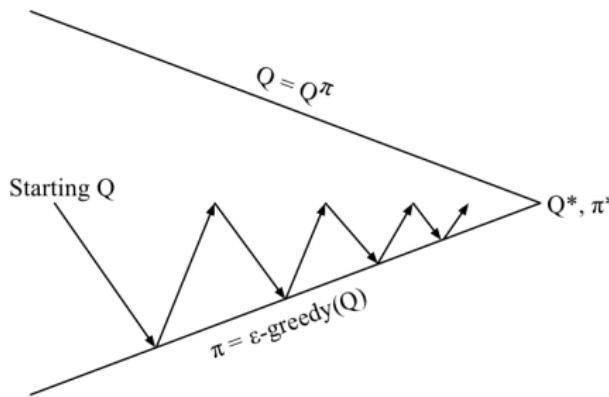
TD(λ)

Model-free
Control

On-Policy
Monte–Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning



Every episode:

- **Policy Evaluation:** Monte–Carlo policy evaluation,
$$Q \approx Q^\pi$$
- **Policy Improvement:** ϵ –greedy policy improvement



Definition

Greedy in the Limit of Infinite Exploration (GLIE)

- All state-action pairs are explored **infinitely** many times

$$\lim_{t \rightarrow \infty} N_k(s, a) = \infty$$

- The policy **converges** on a **greedy** policy

$$\lim_{t \rightarrow \infty} \pi_k(a|s) = \mathbf{1}(a = \arg \max_{a' \in \mathcal{A}} Q_k(s', a'))$$



GLIE Monte–Carlo Control

Marcello
Restelli

Model-free
Prediction

Monte–Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte–Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- Sample k -th episode using π : $\{s_1, a_1, r_2, \dots, s_T\} \sim \pi$
- For each state s_t and action a_t in the episode,

$$N(s_t, a_t) \leftarrow N(s_t, a_t) + 1$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \frac{1}{N(s_t, a_t)}(v_t - Q(s_t, a_t))$$

- **Improve policy** based on new action–value function

$$\epsilon \leftarrow \frac{1}{k}$$

$$\pi \leftarrow \epsilon\text{-greedy}(Q)$$

Theorem

*GLIE Monte–Carlo control **converges** to the **optimal** action–value function, $Q(s, a) \rightarrow Q^*(s, a)$*



Relevant Time Scales

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $\text{TD}(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- There are three main time scales
 - 1 **Behavioral** time scale $\frac{1}{1-\gamma}$ (**discount factor**)
 - 2 **Sampling** in the estimation of the Q -function α (**learning rate**)
 - 3 **Exploration** ϵ (e.g., for ϵ -greedy strategy)
- $1 - \gamma \gg \alpha \gg \epsilon$
- **Initially** $1 - \gamma \approx \alpha \approx \epsilon$ is possible
- Then decrease ϵ **faster** than α
- **Practically**, you can choose number of trials $M < \infty$ and set $\alpha \sim 1 - \frac{m}{M}$ and $\epsilon \sim (1 - \frac{m}{M})^2$, $m = 1, \dots, M$
- In some cases, γ should be **initialized to low values** and then gradually moved towards its correct value



MC vs TD Control

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

$\text{TD}(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- Temporal-Difference (TD) learning has several **advantages** over Monte-Carlo (MC)
 - Lower Variance
 - Online
 - Incomplete sequences
- Natural idea: use **TD** instead of MC in our **control loop**
 - Apply TD to $Q(s, a)$
 - Use ϵ -greedy policy improvement
 - Update every time-step



On-Policy Control with SARSA

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

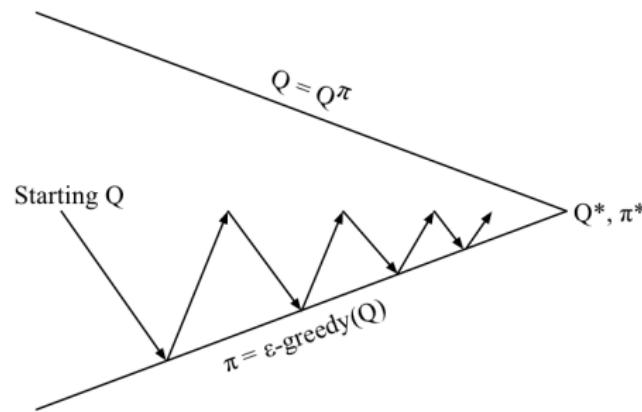
Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$$



Every time-step:

- **Policy Evaluation:** **SARSA**, $Q \approx Q^\pi$
- **Policy Improvement:** ϵ -greedy policy improvement



SARSA Algorithm for On-Policy Control

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Initialize $Q(s, a)$ arbitrarily

loop

 Initialize s

 Choose a from s using policy derived from Q (e.g.,
 ϵ -greedy)

repeat

 Take action a , observe r, s'

 Choose a' from s' using policy derived from Q (e.g.,
 ϵ -greedy)

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$$

$s \leftarrow s'; a \leftarrow a'$;

until s is terminal

end loop



Convergence of SARSA

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Theorem

*SARSA converges to the optimal action-value function,
 $Q(s, a) \rightarrow Q^*(s, a)$, under the following conditions:*

- GLIE sequence of policies $\pi_t(s, a)$
- Robbins–Monro sequence of step-sizes α_t

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$



SARSA Example

Windy Gridworld

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

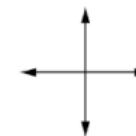
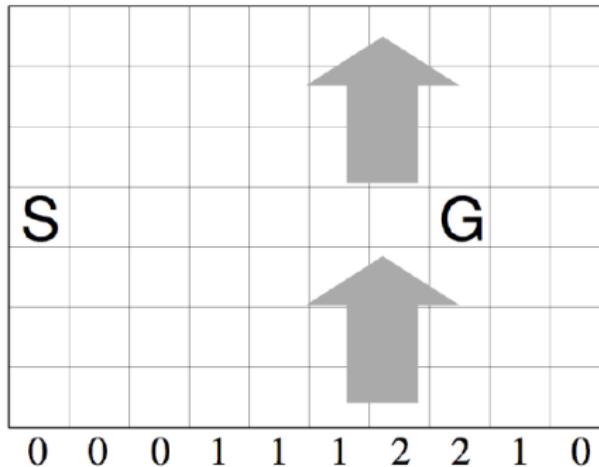
TD(λ)

Model-free
Control

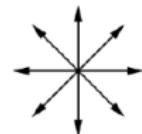
On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning



standard
moves



king's
moves

undiscounted, episodic, reward = -1 until goal



SARSA Example

Results in the Windy Gridworld

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

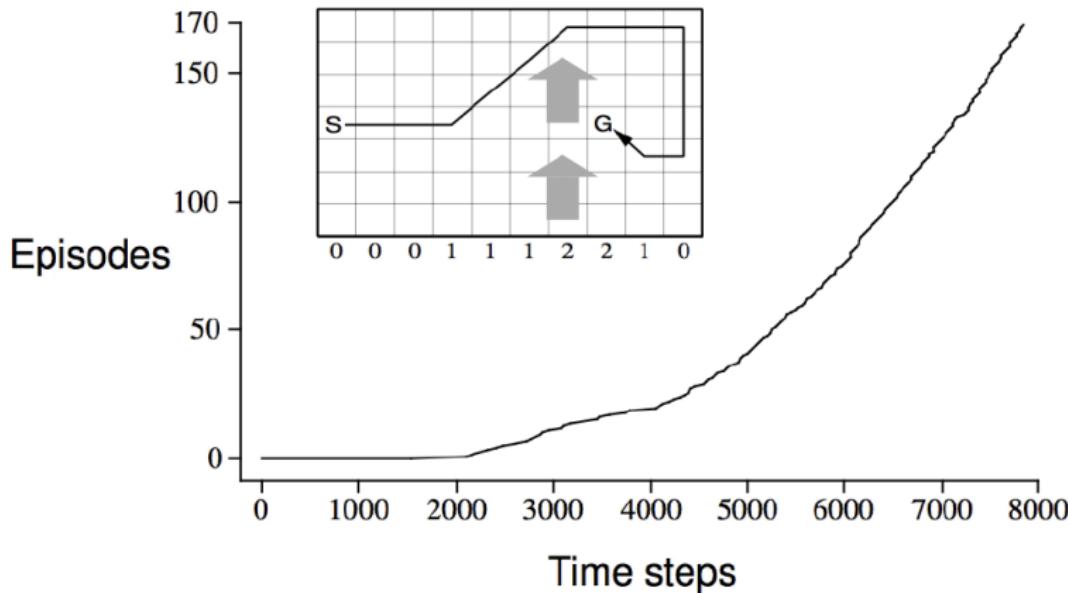
$TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning



SARSA with Eligibility Traces

Marcello
Restelli

Model-free
Prediction

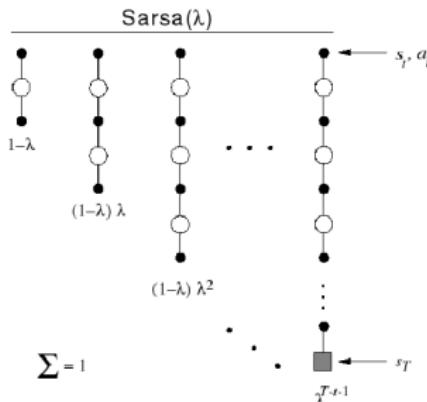
Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $\text{TD}(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning



- **Forward view:** update action-value $Q(s, a)$ to λ -return v_t^λ
- **Backward view:** use eligibility traces for state-action pairs

$$e_t(s, a) = \gamma \lambda e_{t-1}(s, a) + \mathbf{1}(s_t, a_t = s, a)$$



SARSA(λ) Algorithm

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Initialize $Q(s, a)$ arbitrarily

loop

$e(s, a) = 0$, for all s, a

Initialize s, a

repeat

Take action a , observe r, s'

Choose a' from s' using policy derived from Q (e.g., ϵ -greedy)

$\delta \leftarrow r + \gamma Q(s', a') - Q(s, a)$

$e(s, a) \leftarrow e(s, a) + 1$

for all s, a **do**

$Q(s, a) \leftarrow Q(s, a) + \alpha \delta e(s, a)$

$e(s, a) \leftarrow \gamma \lambda e(s, a)$

end for

$s \leftarrow s'; a \leftarrow a'$;

until s is terminal

end loop



SARSA(λ) Gridworld Example

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

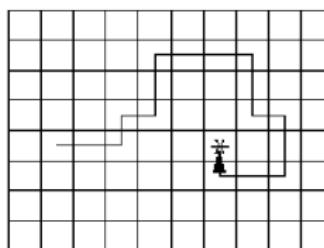
Model-free
Control

On-Policy
Monte-Carlo Control

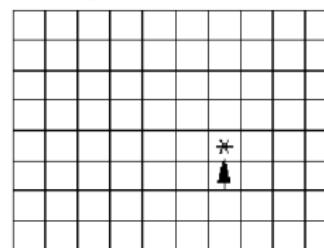
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

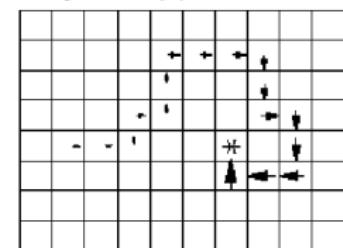
Path taken



Action values increased
by one-step Sarsa



Action values increased
by Sarsa(λ) with $\lambda=0.9$





Off-Policy Learning

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Learn about **target policy** $\pi(a|s)$
- While following **behavior policy** $\bar{\pi}(a|s)$
- Why is this important?
 - Learn from **observing** humans or other agents
 - **Re-use** experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$
- Learn about **optimal** policy while following **exploratory** policy
- Learn about **multiple** policies while following **one** policy



Importance Sampling

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Estimate the expectation of a **different** distribution w.r.t. the distribution used to **draw samples**

$$\begin{aligned}\mathbb{E}_{x \sim P}[f(x)] &= \sum P(x)f(x) \\ &= \sum Q(x)\frac{P(x)}{Q(x)}f(x) \\ &= \mathbb{E}_{x \sim Q}\left[\frac{P(x)}{Q(x)}f(x)\right]\end{aligned}$$



Importance Sampling for Off-Policy Monte-Carlo

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

TD(λ)

Model-free Control

On-Policy Monte-Carlo Control

On-Policy Temporal-Difference Learning

Off-Policy Learning

- Use returns **generated** from $\bar{\pi}$ to **evaluate** π
- Weight return v_t according to **similarity** between policies
- Multiply **importance sampling corrections** along whole episode

$$v_t^\mu = \frac{\pi(a_t|s_t)}{\bar{\pi}(a_t|s_t)} \frac{\pi(a_{t+1}|s_{t+1})}{\bar{\pi}(a_{t+1}|s_{t+1})} \cdots \frac{\pi(a_T|s_T)}{\bar{\pi}(a_T|s_T)} v_t$$

- Update value towards **corrected** return

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(v_t - Q(s_t, a_t))$$

- **Cannot use** if $\bar{\pi}$ is zero where π is non-zero
- Importance sampling can dramatically **increase variance**



Importance Sampling for Off-Policy Monte-Carlo Derivation

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

TD(λ)

Model-free Control

On-Policy Monte-Carlo Control

On-Policy Temporal-Difference Learning

Off-Policy Learning

Off-policy MC is derived from **expected return**:

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E}_\pi[v_t | s_t = s, a_t = a] \\ &= \sum \mathbb{P}[s_1, a_1, r_2, \dots, s_T] v_t \\ &= \sum \mathbb{P}[s_1] \left(\prod_{t=1}^T \bar{\pi}(s_t, a_t) P(s_t | s_{t-1}, a_{t-1}) \frac{\pi(s_t | a_t)}{\bar{\pi}(s_t, a_t)} \right) v_t \\ &= \mathbb{E}_{\bar{\pi}} \left[\prod_{t=1}^T \frac{\pi(s_t, a_t)}{\bar{\pi}(s_t, a_t)} v_t | s_t = s, a_t = a \right] \end{aligned}$$



Off-Policy MC Control

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Initialize, for all $s \in \mathcal{S}, a \in \mathcal{A}$:

$$Q(s, a) \leftarrow \text{arbitrary}$$

$$N(s, a) \leftarrow 0$$

$$D(s, a) \leftarrow 0$$

$\pi \leftarrow$ an arbitrary deterministic policy

loop

Using a policy $\bar{\pi}$, generate an episode

$$s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T, s_T$$

$\tau \leftarrow$ latest time at which $a_\tau \neq \pi(s_\tau)$

for all pair s, a appearing in the episode after τ **do**

$t \leftarrow$ the time of first occurrence (after τ) of s, a

$$w \leftarrow \prod_{k=t+1}^{T-1} \frac{1}{\bar{\pi}(s_k, a_k)}$$

$$N(s, a) \leftarrow N(s, a) + w R_t$$

$$D(s, a) \leftarrow D(s, a) + w$$

$$Q(s, a) \leftarrow \frac{N(s, a)}{D(s, a)}$$

end for

for all $s \in \mathcal{S}$ **do**

$$\pi(s) \leftarrow \arg \max_{a \in \mathcal{A}} Q(s, a)$$

end for

end loop



Importance Sampling for Off-Policy SARSA

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning
Comparison between
MC and TD
 $\text{TD}(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning
Off-Policy Learning

- Use **TD targets** generated from π to evaluate $\bar{\pi}$
- **Weight** TD target $r + \gamma Q(s', a')$ according to **similarity** between policies
- Only need a **single** importance sampling correction

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left(r_{t+1} + \gamma \frac{\pi(a|s)}{\bar{\pi}(a|s)} Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t) \right)$$

- Much **lower variance** than Monte-Carlo importance sampling
- Policies only need to be similar over a **single step**



Importance Sampling for Off-Policy SARSA

Bellman expectation equation

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Off-Policy SARSA comes from Bellman expectation equation for $Q^\pi(s, a)$

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E}_\pi [r_{t+1} + \gamma Q^\pi(s_{t+1}, a_{t+1}) | s_t = s, a_t = a] \\ &= R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) \sum_{a' \in \mathcal{A}} \pi(a'|s') Q^\pi(s', a') \\ &= R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) \sum_{a' \in \mathcal{A}} \bar{\pi}(a'|s') \frac{\pi(a'|s')}{\bar{\pi}(a'|s')} Q^\pi(s', a') \\ &= \mathbb{E}_\mu \left[r_{t+1} + \gamma \frac{\pi(a|s)}{\bar{\pi}(a|s)} Q^\pi(s_{t+1}, a_{t+1}) | s_t = s, a_t = a \right] \end{aligned}$$



Off-Policy Control with Q -learning

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD
 $TD(\lambda)$

Model-free
Control

On-Policy
Monte-Carlo Control
On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- Learn about **optimal policy** $\pi = \pi^*$
- From experience sampled from **behavior policy** $\bar{\pi}$
- Estimate $Q(s, a) \approx Q^*(s, a)$
- Behavior policy **can depend on** $Q(s, a)$
 - e.g., $\bar{\pi}$ could be ϵ -greedy with respect to $Q(s, a)$
 - As $Q(s, a) \rightarrow Q^*(s, a)$, behavior policy $\bar{\pi}$ **improves**



Q-learning Algorithm

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a' \in \mathcal{A}} Q(s', a') - Q(s, a))$$

Initialize $Q(s, a)$ arbitrarily

loop

 Initialize s

repeat

 Choose a from s using policy derived from Q (e.g.,
 ϵ -greedy)

 Take action a , observe r, s'

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$
$$s \leftarrow s';$$

until s is terminal

end loop



SARSA vs Q-learning

Cliffwalking

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

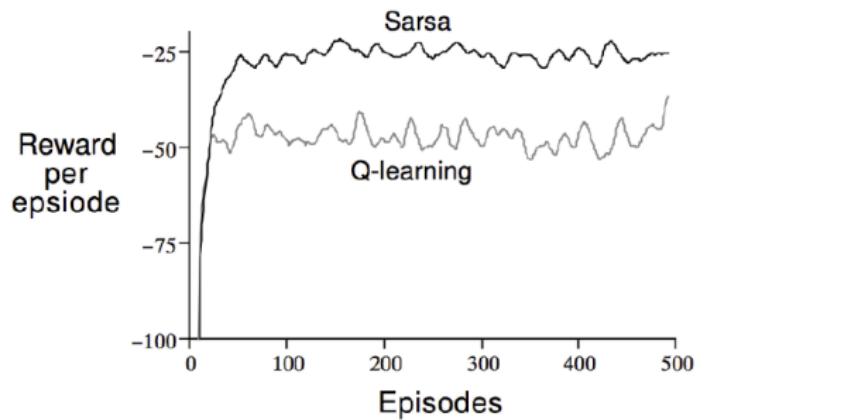
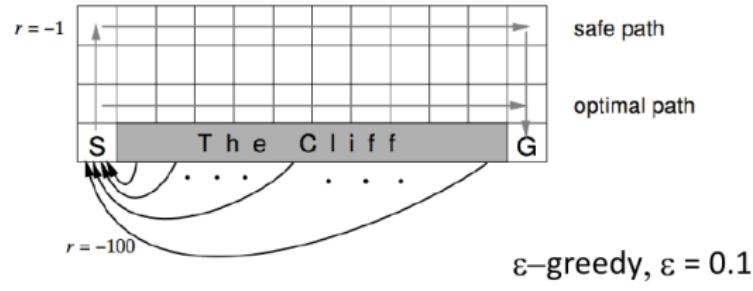
Comparison between MC and TD

$\text{TD}(\lambda)$

Model-free Control

On-Policy Monte-Carlo Control
On-Policy Temporal-Difference Learning

Off-Policy Learning





Q-learning vs SARSA

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- SARSA: $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$
on-policy
- Q-learning:
 $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
off-policy
- In the cliff-walking task:
 - Q-learning: learns **optimal policy** along edge
 - SARSA: learns a **safe non-optimal policy** away from edge
- ϵ -greedy algorithm
 - For $\epsilon \neq 0$ SARSA performs **better online**
 - For $\epsilon \rightarrow 0$ gradually, **both converge to optimal**



Q-Learning, the TD(λ) way

Marcello Restelli

Model-free Prediction

Monte-Carlo Reinforcement Learning

Temporal Difference Reinforcement Learning

Comparison between MC and TD

TD(λ)

Model-free Control

On-Policy

Monte-Carlo Control

On-Policy

Temporal-Difference Learning

Off-Policy Learning

- How we can extend this to Q-learning?
- If you mark every state action pair as eligible, you backup over non-greedy policy

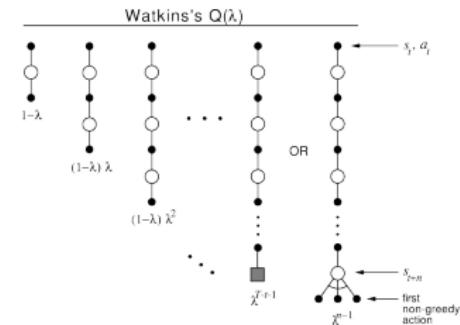
Watkins:

- Zero out eligibility trace after a non-greedy actions
- Do max when backing up at first non-greedy choice

$$e_t(s, a) = \begin{cases} 1 + \gamma \lambda e_{t-1}(s, a) & \text{if } s = s_t, a = a_t, Q_{t-1}(s_t, a_t) = \max_a Q_{t-1}(s_t, a) \\ 0 & \text{if } Q_{t-1}(s_t, a_t) \neq \max_a Q_{t-1}(s_t, a) \\ \gamma \lambda e_{t-1}(s, a) & \text{otherwise} \end{cases}$$

$$Q_{t+1}(s, a) = Q_t(s, a) + \alpha \delta_t e_t(s, a)$$

$$\delta_t = r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q_t(s_t, a_t)$$





Peng's $Q(\lambda)$

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

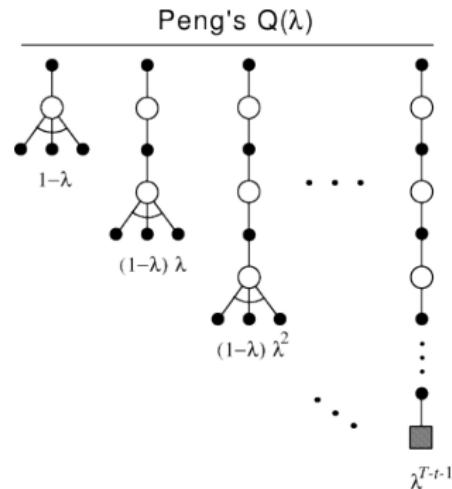
Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

- Disadvantage to Watkins' method:
 - Early in learning, the eligibility trace will be “cut” (zeroed out) frequently resulting in little advantages to traces
- **Peng:**
 - Backup max action except at end
 - Never cut traces
- Disadvantage:
 - Complicated to implement





Relationship Between DP and TD

Marcello
Restelli

Model-free
Prediction

Monte-Carlo
Reinforcement
Learning

Temporal Difference
Reinforcement
Learning

Comparison between
MC and TD

TD(λ)

Model-free
Control

On-Policy
Monte-Carlo Control

On-Policy
Temporal-Difference
Learning

Off-Policy Learning

Full Backup (DP)	Sample backup (TD)
Iterative Policy Evaluation $V(s) \leftarrow \mathbb{E}_\pi[r + \gamma V(s') s]$	TD Learning $V(s) \xleftarrow{\alpha} r + \gamma V(s')$
Q -Policy Iteration $Q(s, a) \leftarrow \mathbb{E}_\pi[r + \gamma Q(s', a') s, a]$	SARSA $Q(s, a) \xleftarrow{\alpha} r + \gamma Q(s', a')$
Q -Value Iteration $Q(s, a) \leftarrow \mathbb{E}_\pi[r + \gamma \max_{a' \in \mathcal{A}} Q(s', a') s, a]$	Q -learning $Q(s, a) \xleftarrow{\alpha} r + \gamma \max_{a' \in \mathcal{A}} Q(s', a')$