

Reinforcement Learning

Policy Gradients

Alberto Maria Metelli

13th February 2024

Book References

Richard S. Sutton, Andrew G. Barto

Reinforcement Learning: An Introduction (second edition)

Chapter 13

Csaba Szepesvári

Algorithms for Reinforcement Learning

Section 4.4



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Value Functions and Policies

- So far, we have learned **parametric value functions**

$$\begin{aligned} v_{\mathbf{w}}(s) &\approx v_{\pi}(s) \text{ or } v_*(s) \\ q_{\mathbf{w}}(s, a) &\approx q_{\pi}(s, a) \text{ or } q_*(s, a) \end{aligned}$$

where $\mathbf{w} \in \mathbb{R}^m$ is the parameter vector

- And a policy is derived from those (e.g., greedy, ϵ -greedy, Boltzmann, ...)
- We now focus on learning **parametric policies**


$$\pi_{\theta}(a|s) = \Pr(\text{playing action } a \text{ in state } s|\theta)$$

where $\theta \in \mathbb{R}^d$ is the parameter vector



Value Based and Policy-Based Reinforcement Learning

- Value-Based
 - **Learn** value function $q_w(s, a)$
 - **Implicit** policy (e.g., greedy, ϵ -greedy, Boltzmann, ...)
- Policy-Based
 - **No** value function
 - **Learn** policy $\pi_\theta(a|s)$
- Actor-Critic
 - **Learn** value function $q_w(s, a)$
 - **Learn** policy $\pi_\theta(a|s)$




Advantages of Policy-Based RL

- Advantages
 - Better **convergence** properties
 - Effective in **high-dimensional** or **continuous action** spaces
 - When $|\mathcal{A}| = \infty$ computing $\sup_{a \in \mathcal{A}} q(s, a)$ is hard!
 - **Policy subspace** can be chosen according to the **task**
 - Policies might be **simpler** than value functions
 - Can learn **stochastic policies**
 - **Exploration** can be directly controlled
 - Better tackle **partial observability** or **non-Markovianity**
 - Can benefit from **demonstrations**
- Disadvantages
 - Typically converge to a **local** rather than a **global** optimum
 - Evaluating a policy is typically **inefficient** and **high variance**



Example: Aliased Gridworld



- The agent **cannot distinguish** the gray states
- Consider **features**:

$$\mathbf{x}(s) = (\mathbf{1}\{\text{wall up}\}, \mathbf{1}\{\text{wall left}\}, \mathbf{1}\{\text{wall right}\}, \mathbf{1}\{\text{wall down}\})^T$$

- Compare value-based RL, using a **parametrized value function**

$$q_{\mathbf{w}}(s, \cdot) = f(\mathbf{x}(s), \mathbf{w})$$


- To policy-based RL, using a **parameterized policy**

$$\pi_{\boldsymbol{\theta}}(\cdot | s) = g(\mathbf{x}(s), \boldsymbol{\theta})$$

Pictures from (Silver, 2015; Hado van Hasselt, 2015)



Example: Aliased Gridworld




- Under aliasing, an optimal **deterministic** policy will either
 - move left in both gray states
 - move right in both gray states
- Either way, it can get stuck and **never** reach the money
- If we add some **uniform noise** over the actions, we will reach the money sooner or later

Pictures from (Silver, 2015; Hado van Hasselt, 2015)



Example: Aliased Gridworld



- An optimal **stochastic** policy will randomly move left or right in gray states

$$\pi_{\theta}(\text{right} \mid \text{wall up} \wedge \text{wall down}) = 0.5$$

$$\pi_{\theta}(\text{left} \mid \text{wall up} \wedge \text{wall down}) = 0.5$$

- It will reach the goal state in a few steps with high probability
- Policy-based RL can learn the optimal stochastic policy

Pictures from (Silver, 2015; Hado van Hasselt, 2015)



Policy Optimization Objective Function

- **Goal:** given a policy $\pi_{\theta}(a|s)$ with parameters θ , find best $\theta \in \mathbb{R}^d$
- But how do we **measure** the quality of a policy π_{θ} ?
- We need a **scalar objective**: the **expected return**

$$J(\theta) = \mathbb{E}_{S_0 \sim d_0} [v_{\pi_{\theta}}(S_0)] = \mathbb{E} \left[\sum_{t=0}^{+\infty} \gamma^t R_{t+1} | S_0 \sim d_0, \pi_{\theta} \right]$$

where d_0 is the **initial state distribution**



Trajectory View

- If we have **trajectories** of finite length T
- We define the **probability of a trajectory** $\tau = (S_0, A_0, S_1, A_1, \dots, S_{T-1}, A_{T-1}, S_T)$

$$p_{\boldsymbol{\theta}}(\tau) = d_0(S_0) \prod_{t=0}^{T-1} \pi_{\boldsymbol{\theta}}(A_t | S_t) p(S_{t+1} | S_t, A_t)$$

- and the **trajectory return**

$$G(\tau) = \sum_{t=0}^{T-1} \gamma^t r(S_t, A_t)$$

- We can rewrite at the **expected return** as

$$J(\boldsymbol{\theta}) = \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [G(\tau)]$$



Occupancy View

- An alternative way, that works also with **infinite-length** trajectories
- We define the γ -**discounted occupancy** (Sutton et al., 1999)

$$d_{\pi_\theta}(s) = \sum_{t=0}^{+\infty} \gamma^t \Pr(S_t = s | S_0 \sim d_0, \pi_\theta)$$

- It is **not** a distribution as it integrates to $\frac{1}{1-\gamma}$
- If the **stationary distribution** of the Markov chain induced by policy π_θ exists, we have

$$\lim_{\gamma \rightarrow 1^-} (1 - \gamma) d_{\pi_\theta}(s) = \text{stationary distribution}$$

- If the MDP has finite-horizon T , the series stops at time instant $T - 1$
- We can rewrite at the **expected return** as

$$J(\theta) = \mathbb{E}_{\substack{S \sim d_{\pi_\theta} \\ A \sim \pi_\theta(\cdot | S)}} [r(S, A)]$$



Occupancy View

- How to sample from d_{π_θ} ?

$S_0 \sim d_0$

for $t = 0, 1, \dots$ **do**

Toss a coin with $1 - \gamma$ head probability (i.e., a Bernoulli r.v. B_t with $p = 1 - \gamma$)

if head (i.e., $B_t = 1$) **then**

return S_t

end if

$A_t \sim \pi_\theta(\cdot | S_t)$

$S_{t+1} \sim p(\cdot | S_t, A_t)$

end for



Trajectory View and Occupancy View are Equivalent

$$\mathbb{E}_{\tau \sim p_{\theta}} [G(\tau)] = \mathbb{E}_{\substack{S \sim d_{\pi_{\theta}} \\ A \sim \pi_{\theta}(\cdot | S)}} [r(S, A)]$$

Proof.

$$\begin{aligned}\mathbb{E}_{\tau \sim p_{\theta}} [G(\tau)] &= \int_{\tau} p_{\theta}(\tau) G(\tau) d\tau = \int_{\tau} d_0(s_0) \prod_{l=0}^{T-1} \pi_{\theta}(a_l | s_l) P(s_{l+1} | s_l, a_l) \sum_{t=0}^{T-1} \gamma^t r(s_t, a_t) d\tau \\ &= \sum_{t=0}^{T-1} \gamma^t \int_{\tau} d_0(s_0) \prod_{l=0}^{T-1} \pi_{\theta}(a_l | s_l) P(s_{l+1} | s_l, a_l) r(s_t, a_t) d\tau \\ &= \sum_{t=0}^{T-1} \gamma^t \int_{\tau_{0:t}} \underbrace{d_0(s_0) \prod_{l=0}^{t-1} \pi_{\theta}(a_l | s_l) P(s_{l+1} | s_l, a_l)}_{\text{past}} \pi_{\theta}(a_t | s_t) r(s_t, a_t) \\ &\quad \times \int_{\tau_{t+1:T}} \underbrace{P(s_{t+1} | s_t, a_t) \prod_{l=t+1}^{T-1} \pi_{\theta}(a_l | s_l) P(s_{l+1} | s_l, a_l)}_{\text{future}} d\tau\end{aligned}$$

Trajectory View and Transition View are Equivalent

Proof.

$$\begin{aligned} &= \sum_{t=0}^{T-1} \gamma^t \int_{s_t, a_t} \underbrace{\int_{\tau_{0:t-1}} d_0(s_0) \prod_{l=0}^{t-1} \pi_{\theta}(a_l | s_l) P(s_{l+1} | s_l, a_l) d\tau_{0:t-1} \pi_{\theta}(a_t | s_t) r(s_t, a_t)}_{= \Pr(S_t = s_t | \pi_{\theta}, S_0 \sim d_0)} ds_t da_t \\ &= \sum_{t=0}^{T-1} \gamma^t \int_{s_t, a_t} \Pr(S_t = s_t | \pi_{\theta}, S_0 \sim d_0) \pi(a_t | s_t) r(s_t, a_t) ds_t da_t \\ &= \int_{s, a} \underbrace{\sum_{t=0}^{T-1} \gamma^t \Pr(s_t = s | \pi_{\theta}, s_0 \sim d_0) \pi_{\theta}(a | s) r(s, a)}_{d_{\pi_{\theta}}(s)} ds da \\ &= \int_{s, a} d_{\pi_{\theta}}(s) \pi_{\theta}(a | s) r(s, a) ds da = \mathbb{E}_{\substack{S \sim d_{\pi_{\theta}} \\ A \sim \pi_{\theta}(\cdot | S)}} [r(S, A)] \end{aligned}$$



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Policy Gradients


- Policy-based reinforcement learning is a **stochastic optimization** problem
- Find θ that maximizes $J(\theta)$
- Some approaches **do not use gradient**
 - Hill climbing
 - Simplex
 - Genetic algorithms
- Greater **efficiency** often possible using gradient
 - Gradient descent
 - Conjugate gradient
 - Quasi–Newton
- We focus on **gradient descent**, many extensions possible
- And on methods that exploit **sequential structure**



Greedy vs Incremental


- **Greedy** updates

$$\boldsymbol{\theta}_{k+1} \in \arg \max_{\boldsymbol{\theta} \in \mathbb{R}^d} \mathbb{E}_{A \sim \pi_{\boldsymbol{\theta}}} [q_{\pi_{\boldsymbol{\theta}_k}}(s, A)]$$



- Potentially **unstable** learning process with **large policy jumps**
- **Policy Gradient** updates

$$\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k + \alpha \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) \Big|_{\boldsymbol{\theta}=\boldsymbol{\theta}_k}$$



- **Stable** learning process with **smooth policy improvement**



Policy Gradient



- Let $J : \mathbb{R}^d \rightarrow \mathbb{R}$ be any **policy objective function** (Peters and Schaal, 2008)
- Policy gradient algorithms search for a **local maximum** in $J(\theta)$ by **gradient ascent**

$$\theta \leftarrow \theta + \Delta\theta \quad \Delta\theta = \alpha \nabla_{\theta} J(\theta)$$

- where $\nabla_{\theta} J(v\theta)$ is the **policy gradient**

$$\nabla_{\theta} J(\theta) = \begin{pmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_d} \end{pmatrix}$$

- and α is a **step-size** parameter (or **learning rate**)



Policy Gradient Methods

- **Black-Box Approaches**
 - Finite-Difference Methods
- **White-Box Approaches**
 - Likelihood Ratio Methods (vanilla policy gradient, natural policy gradient)



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Computing Gradients by Finite Differences

- **Black-box** approach (Sadegh and Spall, 1998)
- To **evaluate** policy gradient of $\pi_{\theta}(a|s)$ with $\theta \in \mathbb{R}^d$
- For each dimension $k \in \{1, \dots, d\}$
 - Estimate k -th **partial derivative** of objective function w.r.t. θ
 - By **perturbing** θ by small amount ϵ in k -th dimension

$$\frac{\partial J(\theta)}{\partial \theta_k} \approx \frac{J(\theta + \epsilon \mathbf{u}_k) - J(\theta)}{\epsilon}$$

where \mathbf{u}_k is unit vector with 1 in k -th component, 0 elsewhere

- Simple, noisy, inefficient, but sometimes effective
- Works for arbitrary policies, even if policy is **not differentiable**
- Do not need to know the **functional form** of $\pi_{\theta}(a|s)$



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



White–Box approach

- We now compute the gradient **analytically** (Peters and Schaal, 2008)
- Policy $\pi_{\theta}(a|s)$ must be **stochastic** and **differentiable** in θ
- Assume we **know** the gradient $\nabla_{\theta}\pi_{\theta}(a|s)$
- $\nabla_{\theta} \log \pi_{\theta}(a|s)$ is called **score function**
- **log trick** identity:

$$\nabla_{\theta} f(\theta) = f(\theta) \frac{\nabla_{\theta} f(\theta)}{f(\theta)} = f(\theta) \nabla_{\theta} \log f(\theta)$$



Likelihood Ratio Gradient

- We can compute the gradient w.r.t. θ

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} \int_{\mathcal{T}} p_{\theta}(\tau) G(\tau) d\tau \\ &= \int_{\mathcal{T}} \nabla_{\theta} p_{\theta}(\tau) G(\tau) d\tau \\ &= \int_{\mathcal{T}} p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) G(\tau) d\tau \\ &= \mathbb{E}_{\tau \sim p_{\theta}} [\nabla_{\theta} \log p_{\theta}(\tau) G(\tau)]\end{aligned}$$

- We have rewritten the gradient as an **expectation** over trajectory
- so we can **estimate** it from samples!



Likelihood Ratio Gradient

- What about $\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau)$?

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau) &= \nabla_{\boldsymbol{\theta}} \log \left(d_0(S_0) \prod_{t=0}^{T-1} \pi_{\boldsymbol{\theta}}(A_t|S_t) P(S_{t+1}|S_t, A_t) \right) \\ &= \nabla_{\boldsymbol{\theta}} \left(\log d_0(S_0) + \sum_{t=0}^{T-1} \log \pi_{\boldsymbol{\theta}}(A_t|S_t) + \sum_{t=0}^{T-1} \log P(S_{t+1}|S_t, A_t) \right) \\ &= \sum_{t=0}^{T-1} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A_t|S_t)\end{aligned}$$



Example: Softmax Policy

- For **finite action spaces** ($|\mathcal{A}| < \infty$), we can use a **softmax policy**
- Weight actions using **linear combination** of features $\mathbf{x}(s, a)^T \boldsymbol{\theta}$ where $\mathbf{x} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}^d$
- Probability of action is proportional to exponential weight

$$\pi_{\boldsymbol{\theta}}(a|s) = \frac{e^{\mathbf{x}(s, a)^T \boldsymbol{\theta}}}{\int_{a'} e^{\mathbf{x}(s, a')^T \boldsymbol{\theta}} da'} \propto e^{\mathbf{x}(s, a)^T \boldsymbol{\theta}}$$

- The **score function** is

$$\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(a|s) = \mathbf{x}(s, a) - \mathbb{E}_{A \sim \pi_{\boldsymbol{\theta}}(\cdot|s)}[\mathbf{x}(s, A)]$$

- Different representation with state features $\mathbf{x} : \mathcal{S} \rightarrow \mathbb{R}^d$ and action weights $\boldsymbol{\theta} : \mathcal{A} \rightarrow \mathbb{R}^d$

$$\pi_{\boldsymbol{\theta}}(a|s) \propto e^{\mathbf{x}(s)^T \boldsymbol{\theta}(a)}$$



Example: Gaussian Policy

- In **continuous action spaces** $\mathcal{A} = \mathbb{R}$, a Gaussian policy is natural
- **Mean** is a linear combination of state features $\mu_{\theta}(s) = \mathbf{x}(s)^T \theta$, where $\theta : \mathcal{S} \rightarrow \mathbb{R}^d$
- **Variance** may be fixed σ^2 , or can also parameterized
- Policy is a **Gaussian**, $a \sim \mathcal{N}(\mu_{\theta}(s), \sigma)$

$$\pi_{\theta}(a|s) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2} \left(\frac{a - \mu_{\theta}(s)}{\sigma}\right)^2\right)$$

- The **score function** is

$$\nabla_{\theta} \log \pi_{\theta}(a|s) = \frac{(a - \mu_{\theta}(s))\mathbf{x}(s)}{\sigma^2}$$

- Can be extended to multidimensional actions $\mathcal{A} = \mathbb{R}^p$



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



REINFORCE

- Recall the **policy gradient** form

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) &= \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau) G(\tau)] \\ &= \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} \left[\left(\sum_{l=0}^{T-1} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A_l | S_l) \right) \left(\sum_{t=0}^{T-1} \gamma^t r(S_t, A_t) \right) \right]\end{aligned}$$

- Simplest idea is to replace the expectation with the sample mean \rightarrow **REINFORCE** (Williams, 1992)



REINFORCE (Williams, 1992)

Initialize θ arbitrarily

for all iterations $k = 1, \dots, K$ **do**

Sample m trajectories $\tau_i = (S_0^i, A_0^i, S_1^i, A_1^i, \dots, S_{T-1}^i, A_{T-1}^i, S_T^i)$ following π_θ

Compute the REINFORCE gradient estimate

$$\hat{\nabla}_\theta^{\text{RF}} J(\theta) = \frac{1}{m} \sum_{i=1}^m \hat{g}_i$$

where
$$\hat{g}_i = \left(\sum_{l=0}^{T-1} \nabla_\theta \log \pi_\theta(A_l^i | S_l^i) \right) \left(\sum_{t=0}^{T-1} \gamma^t r(S_t^i, A_t^i) \right)$$

Update parameters

$$\theta \leftarrow \theta + \alpha \hat{\nabla}_\theta^{\text{RF}} J(\theta)$$

end for
return θ





REINFORCE: Intuition

- \hat{g}_i are **unbiased** estimates of $\nabla_{\theta} J(\theta)$
- $G(\tau_i)$ measures how good is trajectory τ
- Moving in the direction of \hat{g}_i pushes up the log probability of the trajectory, in proportion to how good it is

$$\hat{g}_i = \nabla_{\theta} \log p_{\theta}(\tau_i) G(\tau_i)$$

- **Interpretation:** uses good trajectories as supervised examples
 - Like maximum likelihood in supervised learning
 - good trajectories are made more likely while bad less
 - Trial and Error approach



Pictures from (Schulman, 2016) and (Levine, 2018)

REINFORCE

- Pros
 - Easy to compute
 - Does not use Markov property!
 - Can be used in **partially observable** MDPs without modification

- Cons

- Use a single Monte Carlo estimate $G(\tau)$
- It has possibly a very **large variance**: grows with T (Papini et al., 2019)

$$\mathbb{V}\text{ar} \left[\widehat{\nabla}_{\boldsymbol{\theta}}^{\text{RF}} J(\boldsymbol{\theta}) \right] \leq O \left(\frac{T\kappa R_{\max}}{m(1-\gamma)^2} \right) \quad \text{where} \quad \|\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(a|s)\|_2 \leq \kappa$$

- **Slow** convergence



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises




G(PO)MDP

- We can **reduce the variance** thanks to the **causality property** (Baxter and Bartlett, 2001)
 - Reward collected in the **past** do not depend on actions played in the **future**

$$\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} \left[\left(\sum_{l=0}^{T-1} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A_l | S_l) \right) \left(\sum_{t=0}^{T-1} \gamma^t r(S_t, A_t) \right) \right]$$

$$= \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} \left[\sum_{t=0}^{T-1} \left(\gamma^t r(S_t, A_t) \sum_{l=0}^{T-1} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A_l | S_l) \right) \right]$$

$$= \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} \left[\sum_{t=0}^{T-1} \left(\gamma^t r(S_t, A_t) \underbrace{\sum_{l=0}^t \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A_l | S_l)}_{\text{past}} \right) \right] + \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} \left[\sum_{t=0}^{T-1} \left(\gamma^t r(S_t, A_t) \mathbb{E} \left[\sum_{l=t+1}^{T-1} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A_l | S_l) | \tau_{0:t} \right] \right) \right]$$


- This is a consequence of the **log trick**

$$\mathbb{E}_{A \sim \pi_{\boldsymbol{\theta}}(\cdot | s)} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A | s)] = \int \nabla_{\boldsymbol{\theta}} \pi_{\boldsymbol{\theta}}(a | s) da = \nabla_{\boldsymbol{\theta}} \int \pi_{\boldsymbol{\theta}}(a | s) da = \nabla_{\boldsymbol{\theta}} 1 = 0$$



G(PO)MDP

Initialize θ arbitrarily

for all iterations $k = 1, \dots, K$ **do**

Sample m trajectories $\tau_i = (S_0^i, A_0^i, S_1^i, A_1^i, \dots, S_{T-1}^i, A_{T-1}^i, S_T^i)$ following π_θ

Compute the G(PO)MDP gradient estimate

$$\widehat{\nabla}_\theta^{\text{G(PO)MDP}} J(\theta) = \frac{1}{m} \sum_{i=1}^m \widehat{g}_i$$

where $\widehat{g}_i = \left(\sum_{t=0}^{T-1} \gamma^t r(S_t^i, A_t^i) \sum_{l=0}^t \nabla_\theta \log \pi_\theta(A_l^i | S_l^i) \right)$

Update parameters

$$\theta \leftarrow \theta + \alpha \widehat{\nabla}_\theta^{\text{G(PO)MDP}} J(\theta)$$

end for
return θ




Example: Short Corridor

- Left and Right actions have **reversed** effect in the central state
- -1 reward in every state $\neq G$
- Start in state S
- **Softmax** policy with features

$$\mathbf{x}(s, \text{right}) = (1, 0)^T$$

$$\mathbf{x}(s, \text{left}) = (0, 1)^T$$




Pictures from (Sutton and Barto, 2018)

Example: Short Corridor - G(PO)MDP

- Left and Right actions have **reversed** effect in the central state
- -1 reward in every state $\neq G$
- Start in state S
- **Softmax** policy with features

$$\mathbf{x}(s, \text{right}) = (1, 0)^T$$

$$\mathbf{x}(s, \text{left}) = (0, 1)^T$$



Pictures from (Sutton and Barto, 2018)

G(PO)MDP

- G(PO)MDP has **smaller** variance w.r.t. REINFORCE: no longer depends on T (Papini et al., 2019)

$$\mathbb{V}\text{ar} \left[\hat{\nabla}_{\boldsymbol{\theta}}^{\text{G(PO)MDP}} J(\boldsymbol{\theta}) \right] \leq O \left(\frac{\kappa R_{\max}}{m(1-\gamma)^3} \right) \quad \text{where} \quad \|\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(a|s)\|_2 \leq \kappa$$

- Still the variance is quite large



Policy Gradient and Baselines

- We can further reduce the variance with a **baseline** $\mathbf{b}(\tau) \in \mathbb{R}^d$

$$J(\boldsymbol{\theta}) = \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau) \odot (G(\tau) - \mathbf{b}(\tau))] \quad \odot = \text{element-wise product}$$



- **Unbiased** if $\mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau) \odot \mathbf{b}(\tau)] = \mathbf{0}$
- Computed to **minimize the variance** of the estimator
 - Scalar vs vectorial
 - Time-independent or time-dependent
- **Optimal** vectorial baseline for REINFORCE (Peters and Schaal, 2008)
 - The optimal one is **time-independent**

$$\mathbf{b}^{\text{RF}} = \frac{\mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [(\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau))^2 G(\tau)]}{\mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [(\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau))^2]}$$

- **Peters** time-dependent vectorial baseline for G(PO)MDP (Peters and Schaal, 2008)



Puck World Example



- **Continuous actions** exert small force on puck
- Puck is rewarded for getting **close to target**
- Target location is **reset** every 30 seconds
- Policy is trained using variant (conjugate) of Monte-Carlo policy gradient

Pictures from (Silver, 2015)



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Convergence Results

- Policy gradient is a **stochastic gradient**

$$\theta \leftarrow \theta + \alpha \hat{\nabla} J(\theta) = \theta + \alpha (\nabla_{\theta} J(\theta) + \text{noise})$$

- $J(\theta)$ is **non-convex**
 - Converges **asymptotically** to a **local minimum** (under some technical assumptions) (Yuan et al., 2021)
- **Large variance** of stochastic gradients (growing with the length of the horizon)
- Possible **insufficient exploration**: naïve stochastic exploration
- **Global convergence** under some specific assumptions (Bhandari and Russo, 2019)



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Going Beyond the Finite-Horizon

- So far, we considered **finite-length** trajectories
- What about **infinite-length** trajectories?

Theorem (Policy Gradient Theorem (Sutton et al., 1999))

For an infinite-horizon MDP, let π_θ be a **stochastic policy differentiable in θ** , the policy gradient is given by:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\substack{S \sim d_{\pi_\theta} \\ A \sim \pi_\theta(\cdot|S)}} [\nabla_\theta \log \pi_\theta(A|S) q_{\pi_\theta}(S, A)]$$



Proof of the Policy Gradient Theorem

Proof.

First of all, we observe:

$$\nabla_{\theta} J(\theta) = \int d_0(s) \nabla_{\theta} v_{\pi_{\theta}}(s) ds$$

Consider the **Bellman equation**:

$$q_{\pi_{\theta}}(s, a) = r(s, a) + \gamma \int p(y|s, a) v_{\pi_{\theta}}(y) dy$$

We derive the **Bellman equation for the gradient**:

$$\begin{aligned}\nabla_{\theta} v_{\pi_{\theta}}(s) &= \nabla_{\theta} \left(\int \pi_{\theta}(a|s) q_{\pi_{\theta}}(s, a) da \right) \\ &= \int \nabla_{\theta} \pi_{\theta}(a|s) q_{\pi_{\theta}}(s, a) da + \int \pi_{\theta}(a|s) \nabla_{\theta} q_{\pi_{\theta}}(s, a) da \\ &= \int \nabla_{\theta} \pi_{\theta}(a|s) q_{\pi_{\theta}}(s, a) da + \underbrace{\gamma \int \pi_{\theta}(a|s) \int p(y|s, a) \nabla_{\theta} v_{\pi_{\theta}}(y) dy da}_{=f(s)}\end{aligned}$$

Proof of the Policy Gradient Theorem

Proof.

Multiply by $d_{\pi_\theta}(s)$ and integrate over the states

$$\begin{aligned} \int d_{\pi_\theta}(s) f(s) ds &= \int d_{\pi_\theta}(s) \gamma \int \pi_\theta(a|s) \int p(y|s, a) \nabla_\theta v_{\pi_\theta}(y) dy da ds \\ &= \int \sum_{t=0}^{+\infty} \gamma^t \Pr(S_t = s | \pi_\theta, S_0 \sim d_0) \int \pi_\theta(a|s) \int p(y|s, a) \nabla_\theta v_{\pi_\theta}(y) dy da ds \\ &= \int \left(\sum_{t=0}^{+\infty} \gamma^{t+1} \Pr(S_{t+1} = y | \pi_\theta, S_0 \sim d_0) \right) \nabla_\theta v_{\pi_\theta}(y) dy \\ &= \int \left(\sum_{t=0}^{+\infty} \gamma^{t+1} \Pr(S_{t+1} = y | \pi_\theta, S_0 \sim d_0) + d_0(y) - d_0(y) \right) \nabla_\theta v_{\pi_\theta}(y) dy \\ &= \int d_{\pi_\theta}(y) \nabla_\theta v_{\pi_\theta}(y) dy - \int d_0(y) \nabla_\theta v_{\pi_\theta}(y) dy \end{aligned}$$

Integrating the gradient of the value function:

$$\int \cancel{d_{\pi_\theta}(s) \nabla_\theta v_{\pi_\theta}(s)} ds = \int d_{\pi_\theta}(s) \int \nabla_\theta \pi_\theta(a|s) q_{\pi_\theta}(s, a) da ds + \int \cancel{d_{\pi_\theta}(y) \nabla_\theta v_{\pi_\theta}(y)} dy - \underbrace{\int d_0(y) \nabla_\theta v_{\pi_\theta}(y) dy}_{\nabla_\theta J(\theta)}$$



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Reducing Variance using Critic

- G(PO)MDP still has a **high variance**
- We use a **critic** to estimate the action–value function

$$q_{\mathbf{w}}(s, a) \approx q_{\pi_{\theta}}(s, a)$$

- Actor–critic algorithms maintain **two** sets of parameters
 - **Critic:** Updates **action–value function** parameters $\mathbf{w} \in \mathbb{R}^n$
 - **Actor:** Updates **policy parameters** $\theta \in \mathbb{R}^d$, in direction suggested by critic
- Actor–critic algorithms follow an **approximate policy gradient** via policy gradient theorem

$$\begin{aligned}\nabla_{\theta} J(\theta) &\approx \mathbb{E}_{\substack{S \sim d_{\pi_{\theta}} \\ A \sim \pi_{\theta}(\cdot | S)}} [\nabla_{\theta} \log \pi_{\theta}(A | S) q_{\mathbf{w}}(S, A)] \\ \Delta \theta &= \alpha \nabla_{\theta} \log \pi_{\theta}(A | S) q_{\mathbf{w}}(S, A)\end{aligned}$$



Estimating the Action–Value Function

- Computing the critic is a **policy evaluation** problem
 - G(PO)MDP is equivalent to estimate $q_w(S_t, A_t)$ with a single MC simulation

$$q_w(S_t, A_t) \equiv G_t = \sum_{l=t}^{T-1} \gamma^{l-t} r(S_l, A_l)$$

- Monte Carlo policy evaluation
- Temporal–Difference learning (TD(0), TD(λ))
- Least–Squares Policy Evaluation



Action–Value Actor–Critic

- Using linear value function approximation $q_{\mathbf{w}}(s, a) = \mathbf{x}(s, a)^T \mathbf{w}$
 - **Critic:** Updates \mathbf{w} by linear semi-gradient TD(0) and learning rate β
 - **Actor:** Updates θ by policy gradient theorem and learning rate α

Initialize $\theta \in \mathbb{R}^n$ and $\mathbf{w} \in \mathbb{R}^d$

loop for each episode

 Initialize S

 Sample $A \sim \pi_{\theta}(\cdot | S)$

 Take action A , observe reward R , and next state S'

loop for each step of the episode

 Sample $A' \sim \pi_{\theta}(\cdot | S')$

$\delta \leftarrow R + \gamma q_{\mathbf{w}}(S', A') - q_{\mathbf{w}}(S, A)$

 Update critic $\mathbf{w} \leftarrow \mathbf{w} + \beta \delta \nabla_{\mathbf{w}} q_{\mathbf{w}}(S, A)$

 Update actor $\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(A | S) q_{\mathbf{w}}(S, A)$

$S \leftarrow S'$

$A \leftarrow A'$

 Take action A , observe reward R , and next state S'

end loop

end loop



Bias in Actor–Critic Algorithms

- $q_w(s, a)$ is a **biased** estimate of $q_{\pi_\theta}(s, a)$
- The update of θ may not follow the gradient of $\nabla_\theta J(\theta)$
- A biased policy gradient may **not** find the right solution
- If we choose action–value function approximation $q_w(s, a)$ **carefully**, we can **avoid** any bias!



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Compatible Function Approximation

Theorem (Compatible Function Approximation Theorem Sutton et al. (1999))

An action-value function $q_{\mathbf{w}}(s, a)$ is **compatible** with the policy space $\pi_{\boldsymbol{\theta}}$ if:

- ① The following identity between gradients hold:

$$\nabla_{\mathbf{w}} q_{\mathbf{w}}(s, a) = \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(a|s) \quad \forall (s, a) \in \mathcal{S} \times \mathcal{A}$$

- ② Value function parameters \mathbf{w} minimize the mean square value error under the γ -discounted occupancy:

$$\mathbf{w} \in \arg \min_{\mathbf{w} \in \mathbb{R}^d} \overline{VE}(\mathbf{w}) := \frac{1}{2} \mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} [(q_{\pi_{\boldsymbol{\theta}}}(S, A) - q_{\mathbf{w}}(S, A))^2]$$

Then, the policy gradient computed replacing $q_{\pi_{\boldsymbol{\theta}}}(s, a)$ with $q_{\mathbf{w}}(s, a)$ is exact, i.e.,

$$\mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S) q_{\mathbf{w}}(S, A)] = \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta})$$



Proof of Compatible Function Approximation Theorem

Proof.

If \mathbf{w} is chosen to **minimize** mean square value error, gradient of $\overline{VE}(\mathbf{w})$ w.r.t. \mathbf{w} must be zero:

$$\mathbf{0} = \nabla_{\mathbf{w}} \overline{VE}(\mathbf{w}) = \mathbb{E}_{\substack{S \sim d_{\pi_{\theta}} \\ A \sim \pi_{\theta}(\cdot|S)}} [(q_{\pi_{\theta}}(S, A) - q_{\mathbf{w}}(S, A)) \nabla_{\mathbf{w}} q_{\mathbf{w}}(S, A)] \quad (\text{condition (2)})$$

$$= \mathbb{E}_{\substack{S \sim d_{\pi_{\theta}} \\ A \sim \pi_{\theta}(\cdot|S)}} [(q_{\pi_{\theta}}(S, A) - q_{\mathbf{w}}(S, A)) \nabla_{\theta} \log \pi_{\theta}(A|S)] \quad (\text{condition (1)})$$

$$\implies \mathbb{E}_{\substack{S \sim d_{\pi_{\theta}} \\ A \sim \pi_{\theta}(\cdot|S)}} [\nabla_{\theta} \log \pi_{\theta}(A|S) q_{\mathbf{w}}(S, A)] = \mathbb{E}_{\substack{S \sim d_{\pi_{\theta}} \\ A \sim \pi_{\theta}(\cdot|S)}} [\nabla_{\theta} \log \pi_{\theta}(A|S) q_{\pi_{\theta}}(S, A)] = \nabla_{\theta} J(\theta)$$

□

- Actually, it is necessary that \mathbf{w} is just a **stationary point** of $\overline{VE}(\mathbf{w})$
- A straightforward choice:

$$q_{\mathbf{w}}(s, a) = \mathbf{w}^T \nabla_{\theta} \log \pi_{\theta}(a|s)$$



Actor-Critic with a Baseline

- Similarly to the trajectory case, we can use a **baseline** $\mathbf{b}(s)$

$$\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = \mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S) \odot (q_{\pi_{\boldsymbol{\theta}}}(S, A) - \mathbf{b}(S))]$$

- This can **reduce variance**, without biasing

$$\mathbb{E}_{A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S) \odot \mathbf{b}(S)] = \mathbf{0}$$

- A **good** (but slightly suboptimal) choice is the state value function $b(s) = v_{\pi_{\boldsymbol{\theta}}}(s)$
- So we can rewrite the policy gradient using the **advantage function** $A_{\pi_{\boldsymbol{\theta}}}(s, a)$

$$A_{\pi_{\boldsymbol{\theta}}}(s, a) = q_{\pi_{\boldsymbol{\theta}}}(s, a) - v_{\pi_{\boldsymbol{\theta}}}(s)$$

$$\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = \mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S) A_{\pi_{\boldsymbol{\theta}}}(S, A)]$$

- The advantage function can notably reduce the variance!



Actor-Critic with a Advantage Function

- We could estimate $v_{\mathbf{v}}(s) \approx v_{\pi}(s)$ and $q_{\mathbf{w}}(s, a) \approx q_{\pi}(s, a)$ **independently**
- $A_{\mathbf{w}, \mathbf{v}}(s, a) = q_{\mathbf{w}}(s, a) - v_{\mathbf{v}}(s)$ is **biased** and **unstable**
- Instead, we consider the TD-error

$$\delta_{\pi_{\boldsymbol{\theta}}}(s, a, s') = r(s, a) + \gamma v_{\pi_{\boldsymbol{\theta}}}(s') - v_{\pi_{\boldsymbol{\theta}}}(s)$$

- $\delta_{\pi_{\boldsymbol{\theta}}}(s, a, s')$ is an **unbiased** estimate for $A_{\pi_{\boldsymbol{\theta}}}(s, a)$

$$\mathbb{E}_{S' \sim p(\cdot | S, A)} [r(S, A) + \gamma v_{\pi_{\boldsymbol{\theta}}}(S') - v_{\pi_{\boldsymbol{\theta}}}(S)] = q_{\pi_{\boldsymbol{\theta}}}(S, A) - v_{\pi_{\boldsymbol{\theta}}}(S)$$

- In practice, we estimate just $v_{\mathbf{v}}(s) \approx v_{\pi_{\boldsymbol{\theta}}}(s)$ and approximate:


$$A_{\pi_{\boldsymbol{\theta}}}(s, a) \approx r(s, a) + \gamma v_{\mathbf{v}}(s') - v_{\mathbf{v}}(s)$$



Example: Short Corridor

- Left and Right actions have **reversed** effect in the central state
- -1 reward in every state $\neq G$
- Start in state S
- Advantage function estimated as:

$$A_{\pi_\theta}(S_t, A_t) \approx G_t - v_v(S_t)$$



Pictures from (Sutton and Barto, 2018)

Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Alternative Policy Gradient Directions

- Gradient ascent algorithms can follow **any** ascent direction $\mathbf{d} \in \mathbb{R}^d$, i.e.

$$\mathbf{d}^T \nabla_{\theta} f(\theta) > 0$$

- The **steepest ascent direction**, i.e., the **vanilla gradient** $\mathbf{d} = \nabla_{\theta} f(\theta)$ yields the **most increase** of $f(\theta)$ per “unit” of change in θ

$$\frac{\nabla_{\theta} f(\theta)}{\|\nabla_{\theta} f(\theta)\|_2} = \lim_{\epsilon \rightarrow 0} \arg \max_{\mathbf{d} \in \mathbb{R}^d: \|\mathbf{d}\|_2 \leq \epsilon} \underbrace{\frac{f(\theta + \mathbf{d}) - f(\theta)}{\epsilon}}_{\approx \text{incremental ratio}}$$

- The “unit” of change in θ is measured in **Euclidean** distance $\|\cdot\|_2$
- Distance should be chosen based on the **manifold** and not based on **coordinates** (Amari, 1998)
- What if we change the **distance**?



Natural Gradient

- In a **Riemannian** space, the distance is defined as (Amari, 1998)

$$d(\boldsymbol{\theta}, \boldsymbol{\theta} + \Delta\boldsymbol{\theta}) = \|\Delta\boldsymbol{\theta}\|_{\mathbf{G}(\boldsymbol{\theta})}^2 := \Delta\boldsymbol{\theta}^\top \mathbf{G}(\boldsymbol{\theta}) \Delta\boldsymbol{\theta}$$

where $\mathbf{G}(\boldsymbol{\theta})$ is the **metric tensor** (positive definite matrix)

- In the Euclidian space we have $\mathbf{G}(\boldsymbol{\theta}) = \text{Id}$
- The **steepest ascent in a Riemannian space** is given by Ollivier et al. (2017)

$$\frac{\mathbf{G}(\boldsymbol{\theta})^{-1} \nabla_{\boldsymbol{\theta}} f(\boldsymbol{\theta})}{\|\nabla_{\boldsymbol{\theta}} f(\boldsymbol{\theta})\|_{\mathbf{G}(\boldsymbol{\theta})^{-1}}} = \lim_{\epsilon \rightarrow 0} \arg \max_{\mathbf{d} \in \mathbb{R}^d : \|\mathbf{d}\|_{\mathbf{G}(\boldsymbol{\theta})} \leq \epsilon} \frac{f(\boldsymbol{\theta} + \mathbf{d}) - f(\boldsymbol{\theta})}{\epsilon}$$

- The corresponding direction is called **natural gradient**

$$\tilde{\nabla}_{\boldsymbol{\theta}} f(\boldsymbol{\theta}) = \mathbf{G}(\boldsymbol{\theta})^{-1} \nabla_{\boldsymbol{\theta}} f(\boldsymbol{\theta})$$

- This is an **ascent direction** as $\mathbf{G}(\boldsymbol{\theta})$ is positive definite
- How to select the metric tensor $\mathbf{G}(\boldsymbol{\theta})$?



Fisher Information Matrix

- A common choice in ML is the **Fisher Information Matrix** (FIM) as metric tensor

$$\mathbf{F}(\boldsymbol{\theta}) = \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [\nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau) \nabla_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau)^T] = -\mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} [\mathcal{H}_{\boldsymbol{\theta}} \log p_{\boldsymbol{\theta}}(\tau)]$$

- The FIM has some remarkable properties
 - It is the Hessian of the **KL-divergence** between two distributions

$$\lim_{\|\Delta\boldsymbol{\theta}\|_2 \rightarrow 0} \mathcal{H}_{\boldsymbol{\theta}} D_{\text{KL}}(p_{\boldsymbol{\theta} + \Delta\boldsymbol{\theta}} \| p_{\boldsymbol{\theta}}) = \mathbf{F}(\boldsymbol{\theta}) \quad D_{\text{KL}}(p_{\boldsymbol{\theta}'} \| p_{\boldsymbol{\theta}}) = \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}'}} \left[\log \frac{p_{\boldsymbol{\theta}'}(\tau)}{p_{\boldsymbol{\theta}}(\tau)} \right]$$

- The second-order **Taylor expansion** of the KL-divergence

$$D_{\text{KL}}(p_{\boldsymbol{\theta} + \Delta\boldsymbol{\theta}}, p_{\boldsymbol{\theta}}) = \frac{1}{2} \Delta\boldsymbol{\theta}^T \mathbf{F}(\boldsymbol{\theta}) \Delta\boldsymbol{\theta} + o(\|\Delta\boldsymbol{\theta}\|^3)$$



Natural Gradient

- A step of **vanilla gradient** $\nabla_{\theta} f(\theta)$ controls the Euclidean distance $\|\theta' - \theta\|_2$

$$\Delta\theta = \alpha \frac{\nabla_{\theta} f(\theta)}{\|\nabla_{\theta} f(\theta)\|_2} \implies D_{\text{KL}}(p_{\theta+\Delta\theta}, p_{\theta}) \simeq \frac{\alpha^2}{2} \frac{\|\nabla_{\theta} f(\theta)\|_{\mathbf{F}(\theta)}}{\|\nabla_{\theta} f(\theta)\|_2}$$

- A step of **natural gradient** $\tilde{\nabla}_{\theta} f(\theta) = \mathbf{F}(\theta)^{-1} \nabla_{\theta} f(\theta)$ controls the KL-divergence $D_{\text{KL}}(p_{\theta'} \| p_{\theta})$

$$\Delta\theta = \alpha \frac{\mathbf{F}(\theta)^{-1} \nabla_{\theta} f(\theta)}{\|\nabla_{\theta} f(\theta)\|_{\mathbf{F}(\theta)^{-1}}} \implies D_{\text{KL}}(p_{\theta+\Delta\theta}, p_{\theta}) \simeq \frac{\alpha^2}{2}$$

- Thus, natural gradient is **invariant** to the parametrization p_{θ}



Example of Invariance to Parametrization

- Two parametrizations:

$$f_1(\theta) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | \theta, \sigma^2)}[r(Y)]$$

$$f_2(\rho) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | 2\rho, \sigma^2)}[r(Y)]$$

- We have that $f_1(\theta) = f_2(\rho)$ when $2\rho = \theta$
- Vanilla gradients

$$\nabla_{\theta} f_1(\theta) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | \theta, \sigma^2)} \left[\frac{Y - \theta}{\sigma^2} r(Y) \right]$$

$$\nabla_{\rho} f_2(\rho) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | 2\rho, \sigma^2)} \left[\frac{2(Y - 2\rho)}{\sigma^2} r(Y) \right]$$

- When $2\rho = \theta$ we have that $2\nabla_{\theta} f_1(\theta) = \nabla_{\rho} f_2(\rho)$
- Suppose we update the corresponding parameters with gradient ascent and learning rate α :

$$\theta' = \theta + \alpha \nabla_{\theta} f_1(\theta)$$

$$\rho' = \rho + \alpha \nabla_{\rho} f_2(\rho)$$

- But $\theta' \neq 2\rho'$ → **parametrization dependence**

$$\theta' = \theta + \alpha \nabla_{\theta} f_1(\theta) = 2\rho + \frac{\alpha}{2} \nabla_{\rho} f_2(\rho) \neq 2\rho + 2\alpha \nabla_{\rho} f_2(\rho) = 2\rho'$$



Example of Invariance to Parametrization

- Fisher information matrices

$$F_1(\theta) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | \theta, \sigma^2)} \left[\left(\frac{Y - \theta}{\sigma^2} \right)^2 \right] = \frac{1}{\sigma^2} \quad F_2(\rho) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | 2\rho, \sigma^2)} \left[\left(\frac{2(Y - 2\rho)}{\sigma^2} \right)^2 \right] = \frac{4}{\sigma^2}$$

- Natural gradients

$$\tilde{\nabla}_\theta f_1(\theta) = F_1(\theta)^{-1} \nabla_\theta f_1(\theta) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | \theta, \sigma^2)} [(Y - \theta)r(Y)]$$

$$\tilde{\nabla}_\rho f_2(\rho) = F_2(\rho)^{-1} \nabla_\rho f_2(\rho) = \mathbb{E}_{Y \sim \mathcal{N}(\cdot | 2\rho, \sigma^2)} \left[\frac{1}{2}(Y - 2\rho)r(Y) \right]$$

- When $2\rho = \theta$ we have that $\frac{1}{2}\nabla_\theta f_1(\theta) = \nabla_\rho f_2(\rho)$
- Suppose we update the corresponding parameters with gradient ascent and learning rate α :

$$\theta' = \theta + \alpha \nabla_\theta f_1(\theta) \quad \rho' = \rho + \alpha \nabla_\rho f_2(\rho)$$

- Now $\theta' = 2\rho' \rightarrow \text{parametrization invariance}$

$$\theta' = \theta + \alpha \nabla_\theta f_1(\theta) = 2\rho + 2\alpha \nabla_\rho f_2(\rho) = 2\rho'$$




Natural Policy Gradient


- In the policy gradient case, the FIM can be simplified as (Kakade, 2001)

$$\mathbf{F}(\boldsymbol{\theta}) = \mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S) \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S)^T]$$

(a) ‘Vanilla’ policy gradients



(b) Natural policy gradients



Pictures from (Peters et al., 2003)



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Natural Actor Critic

- Using **compatible** function approximation (Peters et al., 2005)

$$q_{\mathbf{w}}(s, a) = \mathbf{w}^T \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(a|s)$$

- So the natural policy gradient surprisingly **simplifies**

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) &= \mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S) q_{\mathbf{w}}(A|S)] \\ &= \mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} [\nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S) \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S)^T \mathbf{w}] \\ &= \mathbf{F}(\boldsymbol{\theta}) \mathbf{w} \\ \tilde{\nabla}_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) &= \mathbf{F}(\boldsymbol{\theta})^{-1} \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = \mathbf{w}\end{aligned}$$


- i.e., update actor parameters in direction of critic parameters

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha \mathbf{w}$$


- Episodic versions with time-invariant and time-variant baselines (Peters et al., 2005)




Example




(a) Minimum motor command with splines



(c) Passing through a point with splines



(b) Minimum motor command with motor primitives



(d) Passing through a point with motor primitives

Pictures from (Peters and Schaal, 2008)

- Learn motor plans with policy gradients
 - spline-based trajectory plans
 - nonlinear dynamic motor primitives

- Finite Difference Gradient
- Vanilla Policy Gradient with constant baseline
- Vanilla Policy Gradient with time-variant baseline
- Episodic Natural Actor-Critic with single offset basis functions
- Episodic Natural Actor-Critic with time-variant offset basis functions



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Off-Policy Policy Gradient

- Can we estimate the policy gradient $\nabla_{\theta} J(\theta')$ having samples collected with π_{θ} ?
 - $\pi_{\theta'}$ **target** policy
 - π_{θ} **behavioral** policy
- **Importance weighted** policy gradient

$$\begin{aligned}\nabla_{\theta} J(\theta') &= \mathbb{E}_{\tau \sim p_{\theta}} \left[\frac{p_{\theta'}(\tau)}{p_{\theta}(\tau)} \sum_{t=0}^{T-1} \gamma^t \nabla_{\theta} \log \pi_{\theta'}(A_t | S_t) q_{\pi_{\theta}}(S_t, A_t) \right] \\ &\approx \frac{1}{m} \sum_{i=1}^m \frac{p_{\theta'}(\tau_i)}{p_{\theta}(\tau_i)} \sum_{t=0}^{T-1} \gamma^t \nabla_{\theta} \log \pi_{\theta'}(A_t^i | S_t^i) q_{\pi_{\theta}}(S_t^i, A_t^i) =: \widehat{\nabla}_{\theta} J(\theta'/\theta)\end{aligned}$$

where

$$\frac{p_{\theta'}(\tau)}{p_{\theta}(\tau)} = \frac{d_0(S_0) \prod_{t=0}^{T-1} \pi_{\theta'}(A_t | S_t) P(S_{t+1} | S_t, A_t)}{d_0(S_0) \prod_{t=0}^{T-1} \pi_{\theta}(A_t | S_t) P(S_{t+1} | S_t, A_t)} = \frac{\prod_{t=0}^{T-1} \pi_{\theta'}(A_t | S_t)}{\prod_{t=0}^{T-1} \pi_{\theta}(A_t | S_t)}$$

- The estimator is **unbiased** but...



Curse of Horizon

- Unfortunately the **variance** can explode (Metelli et al., 2018)
- The **variance** grows **exponentially** with the horizon T

$$\mathbb{V}\text{ar} \left[\widehat{\nabla}_{\boldsymbol{\theta}} J(\boldsymbol{\theta}' / \boldsymbol{\theta}) \right] \lesssim \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} \left[\left(\frac{p_{\boldsymbol{\theta}'}(\tau)}{p_{\boldsymbol{\theta}}(\tau)} \right)^2 \right] \leq \sup_{s \in \mathcal{S}} \underbrace{\mathbb{E}_{A \sim \pi_{\boldsymbol{\theta}}(\cdot | s)} \left[\left(\frac{\pi_{\boldsymbol{\theta}'}(A | s)}{\pi_{\boldsymbol{\theta}}(A | s)} \right)^2 \right]}_{\approx \text{policy distance}}^T$$

- Can even be **infinite!** → **Curse of Horizon** (Liu et al., 2020)
- Several **general-purpose** fixes: clipping (Ionides, 2008), normalizations (Kuzborskij et al., 2021), smoothing (Metelli et al., 2021)...
- Some fixes specific for MDPs
 - **Per-decision** importance weighting (Precup et al., 2000)
 - **Stationary** importance weighting (Liu et al., 2018)



Per-Decision Importance Weighting

- Use **causality property** like when deriving G(PO)MDP from REINFORCE

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}') &= \mathbb{E}_{\tau \sim p_{\boldsymbol{\theta}}} \left[\sum_{t=0}^{T-1} \prod_{l=0}^t \frac{\pi_{\boldsymbol{\theta}'}(A_l|S_l)}{\pi_{\boldsymbol{\theta}}(A_l|S_l)} \gamma^t \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}'}(A_t|S_t) q_{\pi_{\boldsymbol{\theta}}}(S_t, A_t) \right] \\ &\approx \frac{1}{m} \sum_{i=1}^m \sum_{t=0}^{T-1} \prod_{l=0}^t \frac{\pi_{\boldsymbol{\theta}'}(A_l^i|S_l^i)}{\pi_{\boldsymbol{\theta}}(A_l^i|S_l^i)} \gamma^t \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}'}(A_t^i|S_t^i) q_{\pi_{\boldsymbol{\theta}}}(S_t^i, A_t^i) =: \hat{\nabla}_{\boldsymbol{\theta}} J^{\text{PD}}(\boldsymbol{\theta}'/\boldsymbol{\theta})\end{aligned}$$

- Usually **smaller variance** than vanilla importance weighting (Metelli et al., 2020)

$$\mathbb{V}\text{ar} \left[\hat{\nabla}_{\boldsymbol{\theta}} J(\boldsymbol{\theta}'/\boldsymbol{\theta}) \right] \lesssim \sum_{t=0}^{T-1} \gamma^{2t} D^t \quad D := \sup_{s \in \mathcal{S}} \mathbb{E}_{A \sim \pi_{\boldsymbol{\theta}}(\cdot|s)} \left[\left(\frac{\pi_{\boldsymbol{\theta}'}(A|s)}{\pi_{\boldsymbol{\theta}}(A|s)} \right)^2 \right]$$



Stationary Importance Weighting

- Exploit the **occupancy view** of policy gradient

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}') &= \mathbb{E}_{\substack{S \sim d_{\pi_{\boldsymbol{\theta}}} \\ A \sim \pi_{\boldsymbol{\theta}}(\cdot|S)}} \left[\frac{d_{\pi_{\boldsymbol{\theta}'}}(S) \pi_{\boldsymbol{\theta}'}(A|S)}{d_{\pi_{\boldsymbol{\theta}}}(S) \pi_{\boldsymbol{\theta}}(A|S)} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}'}(A|S) q_{\pi_{\boldsymbol{\theta}}}(S, A) \right] \\ &\approx \frac{1}{m} \sum_{i=1}^m \frac{d_{\pi_{\boldsymbol{\theta}'}}(S^i) \pi_{\boldsymbol{\theta}'}(A^i|S^i)}{d_{\pi_{\boldsymbol{\theta}}}(S^i) \pi_{\boldsymbol{\theta}}(A^i|S^i)} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}'}(A^i|S^i) q_{\pi_{\boldsymbol{\theta}}}(S^i, A^i) =: \widehat{\nabla}_{\boldsymbol{\theta}}^S J(\boldsymbol{\theta}'/\boldsymbol{\theta})\end{aligned}$$

- The estimator is **unbiased** and mitigates the **curse of horizon**, but...
- Requires samples from the distribution $d_{\pi_{\boldsymbol{\theta}}}$ → possibly inefficient sampling
- The functional form of $d_{\pi_{\boldsymbol{\theta}}}$ is **unknown** (depends on the transition model p) (Liu et al., 2018)



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



Trust Regions

- Samples collected with $\pi_{\theta_{\text{old}}}$ can estimate **accurately** performance of policies π_{θ} “close” to $\pi_{\theta_{\text{old}}}$
- **Stability** can be improved by limiting the updates between subsequent policies
- We use a **divergence** D between policies (e.g., KL-divergence, f -divergence, Wasserstein, ...)
 - **Penalized** approach

$$\max_{\theta \in \mathbb{R}^d} J(\theta) - cD(\pi_{\theta}, \pi_{\theta_{\text{old}}})$$

- **Constraint** approach

$$\begin{aligned} & \max_{\theta \in \mathbb{R}^d} J(\theta) \\ & \text{s.t. } D(\pi_{\theta}, \pi_{\theta_{\text{old}}}) \leq c \end{aligned}$$

where c is a hyperparameter

- Many examples: TRPO (Schulman et al., 2015), PPO (Schulman et al., 2017), POIS (Metelli et al., 2018), ...
- We will see in the next lectures...



Outline

① Policy-Based Reinforcement Learning

② Policy Gradients

Black–Box Approaches

White–Box Approaches

REINFORCE

G(PO)MDP

Convergence Results

③ Policy Gradient Theorem

Actor–Critic Policy Gradient

Compatible Function Approximation

④ Natural Policy Gradient

Natural Actor Critic

⑤ *Off-Policy Policy Gradient

Trust Region Methods

⑥ Exercises



REINFORCE Baseline

Exercise 1

Consider the REINFORCE estimator with baseline:

$$\hat{\nabla}_{\boldsymbol{\theta}}^{\text{RF}} J(\boldsymbol{\theta}) = \frac{1}{m} \sum_{i=1}^m \left(\sum_{t=0}^{T-1} \gamma^t r(S_t^i, A_t^i) - \mathbf{b} \right) \odot \left(\sum_{l=0}^T \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A_l^i | S_l^i) \right)$$

- ① Compute the expression of \mathbf{b} that minimizes the variance of the estimator assuming that:
 - ① the baseline is scalar $b \in \mathbb{R}$;
 - ② the baseline is vectorial $\mathbf{b} \in \mathbb{R}^d$.
- ② Propose an estimator for such a baseline and discuss its biasedness.



References I

- S. Amari. Natural gradient works efficiently in learning. *Neural Comput.*, 10(2):251–276, 1998.
- L. Baird. Residual algorithms: Reinforcement learning with function approximation. In *Machine Learning Proceedings 1995*, pages 30–37. Elsevier, 1995.
- J. Baxter and P. L. Bartlett. Infinite-horizon policy-gradient estimation. *J. Artif. Intell. Res.*, 15:319–350, 2001. doi: 10.1613/jair.806.
- J. Bhandari and D. Russo. Global optimality guarantees for policy gradient methods. *CoRR*, abs/1906.01786, 2019. URL <http://arxiv.org/abs/1906.01786>.
- M. H. Hadouvan Hasselt, Diana Borsa. Reinforcement learning lecture series 2021. <https://deepmind.com/learning-resources/reinforcement-learning-series-2021>, 2015.
- E. L. Ionides. Truncated importance sampling. *Journal of Computational and Graphical Statistics*, 17(2):295–311, 2008.
- S. M. Kakade. A natural policy gradient. In T. G. Dietterich, S. Becker, and Z. Ghahramani, editors, *Advances in Neural Information Processing Systems 14 [Neural Information Processing Systems: Natural and Synthetic, NIPS 2001, December 3-8, 2001, Vancouver, British Columbia, Canada]*, pages 1531–1538. MIT Press, 2001.
- I. Kuzborskij, C. Vernade, A. Gyorgy, and C. Szepesvári. Confident off-policy evaluation and selection through self-normalized importance weighting. In *International Conference on Artificial Intelligence and Statistics*, pages 640–648. PMLR, 2021.
- M. Lee and C. W. Anderson. Convergent reinforcement learning control with neural networks and continuous action search. In *2014 IEEE Symposium on Adaptive Dynamic Programming and Reinforcement Learning (ADPRL)*, pages 1–8. IEEE, 2014.
- S. Levine. Cs 294-112: Deep reinforcement learning, policy gradients. <http://rail.eecs.berkeley.edu/deeprlcourse-fa18/>, 2018.
- Q. Liu, L. Li, Z. Tang, and D. Zhou. Breaking the curse of horizon: Infinite-horizon off-policy estimation. *Advances in Neural Information Processing Systems*, 31, 2018.
- Y. Liu, P.-L. Bacon, and E. Brunskill. Understanding the curse of horizon in off-policy evaluation via conditional importance sampling. In *International Conference on Machine Learning*, pages 6184–6193. PMLR, 2020.
- H. R. Maei. Gradient temporal-difference learning algorithms. 2011.



References II

- H. R. Maei and R. S. Sutton. G_q (λ): A general gradient algorithm for temporal-difference prediction learning with eligibility traces. In *Proceedings of the third conference on artificial general intelligence*, volume 1, pages 91–96. Citeseer, 2010.
- H. R. Maei, C. Szepesvári, S. Bhatnagar, D. Precup, D. Silver, and R. S. Sutton. Convergent temporal-difference learning with arbitrary smooth function approximation. In *Advances in Neural Information Processing Systems 22 (NIPS)*, pages 1204–1212, 2009.
- H. R. Maei, C. Szepesvári, S. Bhatnagar, and R. S. Sutton. Toward off-policy learning control with function approximation. In *International Conference on Machine Learning (ICML)*, 2010.
- A. M. Metelli, M. Papini, F. Faccio, and M. Restelli. Policy optimization via importance sampling. In *Advances in Neural Information Processing Systems 31 (NeurIPS)*, pages 5447–5459, 2018.
- A. M. Metelli, M. Papini, N. Montali, and M. Restelli. Importance sampling techniques for policy optimization. *Journal of Machine Learning Research*, 21(141):1–75, 2020.
- A. M. Metelli, A. Russo, and M. Restelli. Subgaussian and differentiable importance sampling for off-policy evaluation and learning. *Advances in Neural Information Processing Systems*, 34, 2021.
- Y. Ollivier, L. Arnold, A. Auger, and N. Hansen. Information-geometric optimization algorithms: A unifying picture via invariance principles. *J. Mach. Learn. Res.*, 18:18:1–18:65, 2017.
- M. Papini, M. Pirotta, and M. Restelli. Smoothing policies and safe policy gradients. *arXiv preprint arXiv:1905.03231*, 2019.
- J. Peters and S. Schaal. Reinforcement learning of motor skills with policy gradients. *Neural Networks*, 21(4):682–697, 2008. doi: 10.1016/j.neunet.2008.02.003.
- J. Peters, S. Vijayakumar, and S. Schaal. Reinforcement learning for humanoid robotics. In *Proceedings of the third IEEE-RAS international conference on humanoid robots*, pages 1–20, 2003.
- J. Peters, S. Vijayakumar, and S. Schaal. Natural actor-critic. In J. Gama, R. Camacho, P. Brazdil, A. Jorge, and L. Torgo, editors, *Machine Learning: ECML 2005, 16th European Conference on Machine Learning, Porto, Portugal, October 3-7, 2005, Proceedings*, volume 3720 of *Lecture Notes in Computer Science*, pages 280–291. Springer, 2005. doi: 10.1007/11564096__29.



References III

- D. Precup, R. S. Sutton, and S. P. Singh. Eligibility traces for off-policy policy evaluation. In *Proceedings of the Seventeenth International Conference on Machine Learning (ICML 2000)*, pages 759–766. Morgan Kaufmann, 2000.
- P. Sadegh and J. C. Spall. Optimal random perturbations for stochastic approximation using a simultaneous perturbation gradient approximation. *IEEE Transactions on Automatic Control*, 43(10):1480–1484, 1998.
- J. Schulman. Deep reinforcement learning: Policy gradients and q-learning. Technical report, Bay Area Deep Learning School, 2016.
- J. Schulman, S. Levine, P. Abbeel, M. I. Jordan, and P. Moritz. Trust region policy optimization. In *Proceedings of the 32nd International Conference on Machine Learning, ICML*, volume 37, pages 1889–1897. JMLR.org, 2015.
- J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal policy optimization algorithms. *CoRR*, abs/1707.06347, 2017.
- D. Silver. Introduction to reinforcement learning.
<https://deepmind.com/learning-resources/-introduction-reinforcement-learning-david-silver>, 2015.
- R. S. Sutton and A. G. Barto. *Reinforcement learning: An introduction*. MIT press, 2018.
- R. S. Sutton, D. A. McAllester, S. P. Singh, and Y. Mansour. Policy gradient methods for reinforcement learning with function approximation. In *Advances in Neural Information Processing Systems 12 (NIPS)*, pages 1057–1063. The MIT Press, 1999.
- R. S. Sutton, C. Szepesvári, and H. R. Maei. A convergent $O(n)$ temporal-difference algorithm for off-policy learning with linear function approximation. In *Advances in Neural Information Processing Systems 21 (NIPS)*, pages 1609–1616. Curran Associates, Inc., 2008.
- R. S. Sutton, H. R. Maei, D. Precup, S. Bhatnagar, D. Silver, C. Szepesvári, and E. Wiewiora. Fast gradient-descent methods for temporal-difference learning with linear function approximation. In *Proceedings of the 26th Annual International Conference on Machine Learning (ICML)*, volume 382, pages 993–1000, 2009.
- R. J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Mach. Learn.*, 8:229–256, 1992.
doi: 10.1007/BF00992696.
- R. Yuan, R. M. Gower, and A. Lazaric. A general sample complexity analysis of vanilla policy gradient. *arXiv preprint arXiv:2107.11433*, 2021.

