

Reinforcement Learning

Model-Based RL and Planning

Alberto Maria Metelli

14th February 2024

Book References

Richard S. Sutton, Andrew G. Barto

Reinforcement Learning: An Introduction (second edition)

Chapter 8



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Model-Free vs Model-Based RL

- **Model-free RL**

- Learn a **value function** v and/or a **policy** π from experience
- **No** model explicitly represented


- **Model-based RL**

- Learn a **model** from experience
- Plan the value function v and/or policy π




Model-Free vs Model-Based RL

Model-Free RL



Model-Based RL



Why learning a Model?

- Disadvantages
 - First learn a **model**, then construct a **value function** v and/or **policy** π
 - Two sources of **approximation error**
 - Learn the **value function** v directly
 - One source of **approximation error**
- Advantages
 - Model can be learned using **supervised learning** methods
 - Can better represent the **uncertainty**
 - Can be more **sample efficient**



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



What is a Model?

- The true **model** is the **joint** reward-next state distribution $\mathcal{P}(r, s'|s, a)$
- We assume conditional independence between next state and reward and deterministic reward

$$\mathcal{P}(r, s'|s, a) = \delta_{r(s,a)}(r) \cdot p(s'|s, a), \quad \text{where:}$$

- $p(s'|s, a)$ is the **transition model** of the MDP
- $r(s, a)$ is the **reward function** of the MDP
- We assume to **know** the state space \mathcal{S} and the action space \mathcal{A}
- Consider **parametric** representations of both r and p

$$\begin{aligned} p(s'|s, a) &\approx p_{\boldsymbol{\eta}}(s'|s, a) \\ r(s, a) &\approx r_{\boldsymbol{\eta}}(s, a) \end{aligned}$$

where $\boldsymbol{\eta} \in \mathbb{R}^p$ is a vector of real parameters



Model Learning

- **Goal:** estimate p_η and r_η from experience $(S_1, A_1, R_2, S_2, \dots, S_{T-1}, A_{T-1}, R_T, S_T)$
- Can be mapped to a **supervised learning** problem over the dataset

$$\begin{array}{lll} S_1, A_1 & \rightarrow & R_2, S_2 \\ & \vdots & \\ S_{T-1}, A_{T-1} & \rightarrow & R_T, S_T \end{array}$$

- Learning the mapping $s, a \rightarrow r$ is a **regression** problem (as $r(s, a)$ is deterministic)
- Learning the mapping $s, a, \rightarrow s'$ is a **density estimation** problem (as s' is stochastic in general)



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Families of Models

- **Expectation models** (Wan et al., 2019)

- We treat the problem $s, a \rightarrow s'$ as a **regression** as well
- Hoping to learn the **expectation of the next state**

$$\bar{p}_{\eta}(s, a) \approx \mathbb{E}_{S' \sim p(\cdot|s, a)}[S']$$

- Ok with **deterministic environments**

- **Stochastic models**

- We treat the problem $s, a \rightarrow s'$ as a **density estimation**

$$p_{\eta}(s'|s, a) \approx p(s'|s, a)$$

- Ok with also **stochastic environments**



Expectation Models

- Select a **function approximator** $\bar{p}_\eta(s, a)$ (e.g., linear models, neural networks, ...)
- Choose a **loss function** $L : \mathbb{R}^p \rightarrow \mathbb{R}$ (e.g., **mean square error**)
- Find the parameters η minimizing the empirical loss
 - If $\mathcal{S} = \mathbb{R}^k$ and mean square error loss:

$$\hat{\eta} \in \arg \min_{\eta \in \mathbb{R}^p} \frac{1}{T} \sum_{t=1}^T \|S_{t+1} - \bar{p}_\eta(S_t, A_t)\|_2^2$$

- Expectation models can have **disadvantages**
 - The **expected next state** might be not informative
 - Ok if the true value function $v_w(s) = w^T s$ is **linear** in the state

$$\mathbb{E}_{S' \sim p(\cdot | s, a)} [v_w(S')] = w^T \mathbb{E}_{S' \sim p(\cdot | s, a)} [S'] \approx w^T \bar{p}_\eta(s, a)$$

- Ok if the environment is **deterministic**



Stochastic Models

- Select a **density approximator** $p_{\eta}(s'|s, a)$ (e.g., Gaussian processes, deep belief networks, ...)
- Choose a **loss function** $L : \mathbb{R}^p \rightarrow \mathbb{R}$ (e.g., **KL-divergence**)
- Find the parameters η minimizing the empirical loss
 - With KL-divergence loss, the problem becomes a **maximum likelihood problem**:

$$\hat{\eta} \in \arg \max_{\eta \in \mathbb{R}^p} \frac{1}{T} \sum_{t=1}^T \log p_{\eta}(S_{t+1}|S_t, A_t)$$

- Models can be **chained** to predict the n -step future state
- ...but the error accumulates



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Table Lookup Model

- When $|\mathcal{S}| < +\infty$ and $|\mathcal{A}| < +\infty$, we can model p and r as tables (Gheshlaghi Azar et al., 2013)
- Count visits $N(s, a) = \sum_{t=1}^T \mathbb{1}\{(S_t, A_t) = (s, a)\}$ for every state-action pair

$$\widehat{p}(s'|s, a) = \frac{1}{N(s, a)} \sum_{t=1}^T \mathbb{1}\{(S_t, A_t, S_{t+1}) = (s, a, s')\}$$
$$\widehat{r}(s, a) = \frac{1}{N(s, a)} \sum_{t=1}^T \mathbb{1}\{(S_t, A_t) = (s, a)\} R_{t+1}$$


- No explicit parametrization η
- If $r(s, a)$ is deterministic, one sample is enough for $\widehat{r}(s, a)$



AB Example

- Two states A, B
- one action from each state
- stochastic reward
- 8 trajectories of experience (S_1, R_2, \dots)
- Table lookup model from experience

A, 0, B, 0
B, 1
B, 1
B, 1
B, 1
B, 1
B, 1
B, 0



Pictures from (Silver, 2015; Hado van Hasselt, 2015)

Linear Expectation Model

- Given a **feature** representation $\mathbf{x} : \mathcal{S} \rightarrow \mathbb{R}^p$ (Wan et al., 2019)
- We encode every state $s \in \mathcal{S}$ as $\mathbf{x}(s)$
- Expected next state and reward are **linear** functions

$$\mathbf{x}(s') \approx \mathbf{T}(a)\mathbf{x}(s) \quad r(s, a) \approx \boldsymbol{\eta}(a)^T \mathbf{x}(s)$$

where $\mathbf{T}(a) \in \mathbb{R}^{p \times p}$ and $\boldsymbol{\eta}(a) \in \mathbb{R}^p$

- Can be optimized via **gradient descent** over the mean square error loss

$$\min_{\substack{\mathbf{T}(a) \in \mathbb{R}^{p \times p} \\ \boldsymbol{\eta}(a) \in \mathbb{R}^p \\ a \in \mathcal{A}}} \frac{1}{T} \sum_{t=1}^T \|\mathbf{x}(S_{t+1}) - \mathbf{T}(A_t)\mathbf{x}(S_t)\|_2^2 + (R_{t+1} - \boldsymbol{\eta}(A_t)^T \mathbf{x}(S_t))^2$$





- Based on **Gaussian Processes (GPs)** (Deisenroth and Rasmussen, 2011)

$$p(S_{t+1}|S_t, A_t) = \mathcal{N}(S_{t+1}|\mu_t, \Sigma_t)$$

- μ_t and Σ_t are fit from data
- Advantages
 - Reduces model bias (GPs model **any** function)
 - Incorporates **uncertainty** into planning (GPs are Bayesian methods)
 - **Analytical** policy gradient computation
- Disadvantages
 - True stochasticity might not be Gaussian
 - GPs are **computationally expensive**



PILCO - Experimental Results



Pictures from (Deisenroth and Rasmussen, 2011)

Error Bound for Estimated Model

Exercise 1

Let (\hat{p}, \hat{r}) be estimates of the true transition model and reward function (p, r) such that:

$$\sup_{(s,a) \in \mathcal{S} \times \mathcal{A}} D_{\text{TV}}(\hat{p}(\cdot|s,a), p(\cdot|s,a)) \leq \epsilon_p \quad , \quad \|\hat{r} - r\|_\infty \leq \epsilon_r \quad \text{and} \quad \|r\|_\infty \leq R_{\max}.$$

Let v_* be the optimal value function computed with (p, r) and \hat{v}_* be the optimal value function computed with (\hat{p}, \hat{r}) , then it holds that:

$$\|v_* - \hat{v}_*\|_\infty \leq \frac{\epsilon_r}{1-\gamma} + \frac{\gamma R_{\max} \epsilon_p}{(1-\gamma)^2}.$$



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Planning

- Once we have learned a model (p_η, r_η) ...
- ...we can use a **planning** algorithm to solve the MDP
 - Dynamic programming (value iteration, policy iteration)
 - Tree search
 - ...
- **Expensive** if the state/action spaces are large
- **Infeasible** if the state/action spaces are continuous



Sample-Based Planning

- We use the model (p_{η}, r_{η}) to **generate samples only**
- **Simulated experience** from the model

$$S_{t+1} \sim p_{\eta}(\cdot | S_t, A_t)$$

$$R_{t+1} = r_{\eta}(S_t, A_t)$$


- Now, we apply **model-free** RL approaches with the simulated experience
 - MC control
 - SARSA
 - Q-learning



AB Example

- Build the table lookup model from **real** experience
- Apply model-free RL to **simulated** experience

A, 0, B, 0
B, 1
B, 0



- MC learning: $v(A) = 1, v(B) = 0.75$

Pictures from (Silver, 2015; Hado van Hasselt, 2015)



Drawbacks of an Inaccurate Model

- The imperfect model has some drawbacks:
 - The policy produced by planning can be **suboptimal**
 - It is the optimal policy of the **approximate MDP** with (p_η, r_η)
 - Model-based RL is only **as good as the estimated model**
- How to cope with them?
 - When the model is wrong, use **model-free RL**
 - reason about the **uncertainty** on η (e.g., Bayesian approaches)
 - **Combine** model-based and model-free RL

Real experience (true MDP)

$$S_{t+1} \sim p(\cdot | S_t, A_t)$$

$$R_{t+1} = r(S_t, A_t)$$

Simulated experience (approximated MDP)

$$S_{t+1} \sim p_\eta(\cdot | S_t, A_t)$$

$$R_{t+1} = r_\eta(S_t, A_t)$$



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Combining Model-Free and Model-Based RL

- **Model-free RL**
 - Learn a **value function** v and/or a **policy** π from experience
 - **No** model explicitly represented
- **Model-based RL**
 - Learn a **model** from experience
 - Plan the value function v and/or policy π
- **Dyna**
 - Learn a **model** from experience
 - **Learn and plan** the value function v and/or policy π from real and simulated experience



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search


*Open Loop Planning

*Progressive Widening




Dyna Architecture

Dyna Architecture



Tabular Dyna-Q Architecture

- Use **real experience** to (Sutton, 1990, 1991)
 - Learn the $q_*(s, a)$ with tabular Q-learning (**model-free**)
 - Learn the model p and r with table lookup (**model-based**)
- Generate **simulated experience** to
 - Plan for $q_*(s, a)$ with tabular Q-planning



Pictures from (Sutton and Barto, 2018)

Tabular Dyna-Q with Deterministic Environment

- We need an **exploration policy** that selects the action based on q (e.g., ϵ -greedy, Boltzmann, ...)

Initialize q, \hat{p}, \hat{r} arbitrarily

loop for each episode

$S, A \leftarrow$ initial state and action using the exploration policy

loop for each step of episode

 Take action A , observe reward R , and next state S'

$q(S, A) \leftarrow q(S, A) + \alpha [R + \gamma \max_{a' \in \mathcal{A}} q(S', a') - q(S, A)]$

 Update transition model estimate $\hat{p}(S, A) \leftarrow S'$

 Update reward estimate $\hat{r}(S, A) \leftarrow R$

loop for n times

$S \leftarrow$ random previously observed state

$A \leftarrow$ random previously taken action in S

$S' \leftarrow \hat{p}(S, A)$

$R \leftarrow \hat{r}(S, A)$

$q(S, A) \leftarrow q(S, A) + \alpha [R + \gamma \max_{a' \in \mathcal{A}} q(S', a') - q(S, A)]$

end loop


end loop

end loop




Dyna-Q on a Simple Maze

- Four actions: up, down, right, and left
- Reward: 0 everywhere, +1 in the goal
- $\gamma = 0.95$



Pictures from (Sutton and Barto, 2018)

Dyna-Q on a Simple Maze




Pictures from (Sutton and Barto, 2018)




Dyna-Q on a Simple Maze

- Arrow is the greedy action
- No arrow if all actions have the same value

WITHOUT PLANNING ($n=0$)



WITH PLANNING ($n=50$)



Pictures from (Sutton and Barto, 2018)

What if the Model is Wrong?

- Possible problems:
 - The environment may be **stochastic**
 - **Too little** real experience
 - **Bad generalization** of the function approximator
- The suboptimal policy computed by planning can lead to the **correction** of the model errors
- This happens when the model is **optimistic**
- Dyna-Q+ favors exploring **less tried** transitions (s, a, r, s')


$$\hat{r}(s, a) + \kappa \sqrt{\tau(s, a, s')}$$

where $\kappa \geq 0$ and $\tau(s, a)$ is the number of steps elapsed from the last experience of the transition



Change the Environment during Learning


- The changed environment is **harder**



Pictures from (Sutton and Barto, 2018)

Change the Environment during Learning

- The changed environment is **easier**



Pictures from (Sutton and Barto, 2018)



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening




Planning for Action Selection

- So far, we used planning for improving a **value function** defined over the whole state-action space
- We now consider planning for **selecting the next action** to be executed
- Planning **locally** (for the next action) can be easier than planning for the **global** value function
- But, once we played the action, we need to **re-plan** in the next state



Forward Search

- Build a **search tree** with the current state S_t as root
- The MDP model (p, r) is used to generate node successors and rewards
- Do not solve the whole MDP, only the **sub-MDP** starting from the current state → can be easier!




Pictures from (Silver, 2015; Hado van Hasselt, 2015)



Simulation-Based (Forward) Search

- Use forward search with **sample-based planning**
- **Simulate** episodes of experience from the current state
- Do not build the tree, but apply **model-free RL** to find the best action



Pictures from (Silver, 2015; Hado van Hasselt, 2015)

Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Prediction via Monte-Carlo Simulation

- **Goal:** estimate the value function $v_\pi(S_t)$ of the **current state** only
- Given a model (p_η, r_η) and a **simulation policy** π
- Simulate M episodes from the current state S_t

$$\{(S_t^i = \textcolor{red}{S}_t, A_t^i, R_{t+1}^i, \dots, S_{T-1}^i, A_{T-1}^i, R_T^i, S_T^i)\}_{i=1}^M$$

where: $S_{k+1}^i \sim p_\eta(\cdot | S_k^i, A_k^i)$ $R_{k+1}^i = r_\eta(S_k^i, A_k^i)$

- Estimate $v_\pi(S_t)$ with the **Monte-Carlo returns**

$$\hat{v}(S_t) = \frac{1}{M} \sum_{i=1}^M G_t^i \quad \text{where} \quad G_t^i = \sum_{j=t}^{T-1} \gamma^{j-t} R_{j+1}^i$$



Control via Monte-Carlo Simulation

- **Goal:** find the **best action** to be played in the current state
- Given a model (p_η, r_η) and a **simulation policy** π
- For each action $a \in \mathcal{A}$, Simulate M episodes from the current state S_t

$$\{(S_t^i = S_t, A_t^i = a, R_{t+1}^i, \dots, S_{T-1}^i, A_{T-1}^i, R_T^i, S_T^i)\}_{i=1}^N$$

where: $S_{k+1}^i \sim p_\eta(\cdot | S_k^i, A_k^i)$ $R_{k+1}^i = r_\eta(S_k^i, A_k^i)$

- Estimate $q_\pi(S_t, a)$ with the **Monte-Carlo returns**

$$\hat{q}(S_t, a) = \frac{1}{N} \sum_{i=1}^N G_t^i \quad \text{where} \quad G_t^i = \sum_{j=t}^{T-1} \gamma^{j-t} R_{j+1}^i$$

- Select the action maximizing the estimated $\hat{q}(S_t, a)$

$$A_t \in \arg \max_{a \in \mathcal{A}} \hat{q}(S_t, a)$$



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Monte Carlo Tree-Search


- Incrementally build a **search tree** with the visited states, actions, and estimated $q_*(S_t, a)$ from the current state S_t (Browne et al., 2012)
- Repeat the **four** steps (until budget is expired)
 - ① **Selection:** from root reach a node, choosing actions with the **tree policy**
 - ② **Expand:** one (or more) child nodes are added to expand the tree, according to the available actions
 - ③ **Simulation:** a simulation is run from the new node(s), using the **rollout (or default) policy**
 - ④ **Backup:** $\hat{q}(s, a)$ of the path are updated based on the return of the simulation
- Select the action maximizing the estimate in the root node

$$A_t \in \arg \max_{a \in \mathcal{A}} \hat{q}(S_t, a)$$

- Under certain conditions, **converges** to $q_*(S_t, a)$



Monte Carlo Tree-Search



Pictures from (Sutton and Barto, 2018)

Monte Carlo Tree-Search - Policies

- The **tree policy** balances **exploration** and **exploitation**
 - ϵ -greedy on the estimated $\hat{q}(s, a)$
 - UCB policy → UCT (Upper Confidence Tree) (Coquelin and Munos, 2007)

$$\hat{a} \in \arg \max_{a \in \mathcal{A}} \hat{q}(s, a) + \sqrt{\frac{\alpha \log N(s)}{N(s, a)}}$$

- It **improves** during execution
- The **rollout (or default) policy** is fixed
 - e.g., random uniform policy over \mathcal{A}



MCTS Algorithm

```
Create root node ( $S_t$ )
Initialize  $(S_t).N \leftarrow 0$ 
Initialize  $(S_t).V \leftarrow 0$ 
loop within computational budget
     $(S) \leftarrow \text{TREEPOLICY}((S_t))$ 
     $\Delta \leftarrow \text{ROLLOUTPOLICY}(S, H, \gamma)$ 
    BACKUP( $(S), \Delta, \gamma$ )
end loop
return BESTCHILD( $(S_t)$ )
procedure BESTCHILD( $(S)$ )
    for  $(A)$  child of  $(S)$  do
        Compute Q-value  $Q(A) \leftarrow (A).W/(A).N$ 
        Compute the bonus  $B(A) \leftarrow \sqrt{\alpha \log((S).N)/((A).N)}$ 
    end for
    return  $(\arg \max_{a \in \mathcal{A}} \{Q(a) + B(a)\})$ 
end procedure
```

- State-nodes are denoted with (S) where S is the state
 - $(S).V$ is the rollout return from node (S)
 - $(S).N$ is the number of updates to node (S)
 - $(S).R$ is the immediate reward obtained in the transition that has S as next state
- Action-nodes are denoted with (A) where A is the action
 - $(A).W$ is the sum of the returns from node (A)
 - $(A).N$ is the number of updates to node (A)
 - Thus $\hat{q}(S, A) = (A).W/(A).N$



MCTS Algorithm

```
procedure TREEPOLICY( $(S)$ )
  while  $S$  is non-terminal do
     $(A) \leftarrow \text{BESTCHILD}((S))$ 
     $S' \leftarrow p(S, A)$ 
    if  $(S')$  is a child of  $(A)$  then
       $(S').R \leftarrow r(S, A)$ 
       $(S) \leftarrow (S')$ 
    else
      return EXPAND( $(S', A)$ )
    end if
  end while
  return  $(S')$ 
end procedure
```

```
procedure EXPAND( $(S', A)$ )
  Create node  $(S')$  as a child of  $(A)$ 
  Initialize the value  $(S').V \leftarrow 0$ 
  Initialize the count  $(S').N \leftarrow 0$ 
  for action  $A' \in \mathcal{A}$  do
    Create node  $(A')$  as a child of  $(S')$ 
    Initialize  $(A').W \leftarrow 0$ 
    Initialize  $(A').N \leftarrow 0$ 
  end for
  return  $(S')$ 
end procedure
```



MCTS Algorithm

```
procedure ROLLOUTPOLICY( $S, H, \gamma$ )
     $\Delta \leftarrow 0$ 
     $t \leftarrow 0$ 
    while  $S$  is non-terminal and  $t < H$  do
        Choose  $A$  uniformly at random
         $\Delta \leftarrow \gamma^t r(S, A)$ 
         $S \leftarrow p(S, A)$ 
         $t \leftarrow t + 1$ 
    end while
    return  $\Delta$ 
end procedure
```

```
procedure BACKUP( $(S), \Delta, \gamma$ )
     $(S).V \leftarrow \Delta$ 
     $(S).N \leftarrow (S).N + 1$ 
    while  $(S)$  is not root do
         $\Delta \leftarrow (S).R + \gamma\Delta$ 
        Get  $(A)$  parent of  $(S)$ 
         $(A).N \leftarrow (A).N + 1$ 
         $(A).W \leftarrow (A).W + \Delta$ 
        Get  $(S)$  parent of  $(A)$ 
         $(S).N \leftarrow (S).N + 1$ 
    end while
end procedure
```




Example

- 2 actions
- Reward in terminal state only
- **Greedy** tree policy
- **Random** rollout policy




Example - 1



Pictures from (Silver, 2015; Hado van Hasselt, 2015)




Example - 2



Pictures from (Silver, 2015; Hado van Hasselt, 2015)




Example - 3



Pictures from (Silver, 2015; Hado van Hasselt, 2015)




Example - 4



Pictures from (Silver, 2015; Hado van Hasselt, 2015)

Example - 5



Pictures from (Silver, 2015; Hado van Hasselt, 2015)



Advantages of Monte Carlo Tree-Search

- States are evaluated **dynamically**, unlike DP
- Uses sampling to cope with large state spaces
- Works for **black-box** models
- Computationally efficient, anytime, parallelizable
- Can be applied to **games**



Drawbacks of Monte Carlo Tree-Search

- Actions must be **finite** ($|\mathcal{A}| < +\infty$), otherwise we keep exploring actions
- Environment **can be stochastic and continuous**, but the number of next states must be finite, otherwise infinite branching factor (Jonsson et al., 2020)

$$|\{s' \in \mathcal{S} : p(s'|s, a) > 0\}| \leq B < +\infty \quad \forall (s, a) \in \mathcal{S} \times \mathcal{A}$$

- How to cope with these problems?



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Open Loop Planning

- Plan over the **sequence of actions** as opposed to plan over the **policies** (mapping from states to actions) (Bubeck and Munos, 2010)
- We optimize the **open-loop** objective

$$q_{OL}(S_t, a) = \max_{a_{0:\infty} \in \mathcal{A}^\infty} q(S_t, a_{1:\infty}) = \mathbb{E} \left[\sum_{t=0}^{+\infty} \gamma^t r(S_t, \textcolor{red}{a}_t) | a_0 = a \right]$$

where $a_{0:\infty} = (a, a_1, \dots)$ are selected in advance

- Instead, RL optimizes the **closed-loop** objective

$$q_*(S_t, a) = \max_{\pi: \mathcal{S} \rightarrow \mathcal{A}} q_\pi(S_t, a) = \mathbb{E} \left[\sum_{t=0}^{+\infty} \gamma^t r(S_t, A_t) | A_0 = a, A_t = \pi(S_t) \right]$$



Open Loop Planning

- **Computationally intensive:** search in the space of sequences of length H : $|\mathcal{A}|^H$ sequences
- **Optimal** for deterministic environments
- Can be used with **stochastic environments** but ...
 - ... forced to play the same action at time t , regardless the state...
 - ... so **suboptimal** for stochastic environments

$$q_{\text{OL}}(S_t, a) \leq q_*(S_t, a)$$

- Still requires **finite** actions
- Using upper confidence bounds → **Open Loop Optimistic Planning** (Bubeck and Munos, 2010)



Outline

① Model-Free and Model-Based RL

② Model Learning

Families of Models

Examples of Model Approximators

③ Sample-Based Planning

④ Integrated Architectures

Dyna

⑤ Simulation-Based Search

Prediction and Control via Monte-Carlo Simulation

Monte Carlo Tree-Search

*Open Loop Planning

*Progressive Widening



Progressive Widening

- If we have **continuous actions** ($|\mathcal{A}| = +\infty$), we need to **discretize**
- **Fixed discretization:** choose a finite set of k actions

$$\mathcal{A}_k = \{a_1, \dots, a_n\} \subset \mathcal{A}$$

- Never vanishing approximation error
- Can be optimally selected with environment regularities are present
- **Progressive Widening (PW):** adapt the discretization through time (Chaslot et al., 2008)

$$k(t) = \lceil Ct^\alpha \rceil \quad \text{for some } \alpha \in (0, 1) \quad \mathcal{A}_{k(t)} = \{a_1, \dots, a_{k(t)}\} \subset \mathcal{A}$$

- The action set **grows** though time
- PW does not work in stochastic environments with **infinite possible next states** → **Double Progressive Widening** (Couëtoux et al., 2011)



Error Bound for Open Loop Planning

Exercise 2

Consider the optimal value function:

$$v_*(s) = \max_{\pi: \mathcal{S} \rightarrow \mathcal{A}} v_\pi(s) = \mathbb{E} \left[\sum_{t=0}^{+\infty} \gamma^t r(S_t, A_t) | S_0 = s \right]$$

and the open-loop value function:

$$v_{OL}(s) = \max_{a_{0:\infty} \in \mathcal{A}^\infty} \mathbb{E} \left[\sum_{t=0}^{+\infty} \gamma^t r(S_t, a_t) | S_0 = s \right]$$

where $a_{0:\infty} = (a_0, a_1, \dots)$. Prove that:

$$\|v_* - v_{OL}\|_\infty \leq \frac{2\gamma R_{\max}}{(1-\gamma)^2} \left(1 - \min_{s,a \in \mathcal{S} \times \mathcal{A}} \max_{s' \in \mathcal{S}} p(s'|s,a) \right)$$

Note that when p is deterministic, we have **zero error!**



References I

- C. B. Browne, E. Powley, D. Whitehouse, S. M. Lucas, P. I. Cowling, P. Rohlfshagen, S. Tavener, D. Perez, S. Samothrakis, and S. Colton. A survey of monte carlo tree search methods. *IEEE Transactions on Computational Intelligence and AI in games*, 4(1):1–43, 2012.
- S. Bubeck and R. Munos. Open loop optimistic planning. In *COLT 2010-The 23rd Conference on Learning Theory*, 2010.
- G. M. J. Chaslot, M. H. Winands, H. J. v. d. Herik, J. W. Uiterwijk, and B. Bouzy. Progressive strategies for monte-carlo tree search. *New Mathematics and Natural Computation*, 4(03):343–357, 2008.
- P.-A. Coquelin and R. Munos. Bandit algorithms for tree search. In *Uncertainty in Artificial Intelligence*, 2007.
- A. Couëtoux, J.-B. Hoock, N. Sokolovska, O. Teytaud, and N. Bonnard. Continuous upper confidence trees. In *International Conference on Learning and Intelligent Optimization*, pages 433–445. Springer, 2011.
- M. P. Deisenroth and C. E. Rasmussen. PILCO: A model-based and data-efficient approach to policy search. In L. Getoor and T. Scheffer, editors, *Proceedings of the 28th International Conference on Machine Learning, ICML 2011, Bellevue, Washington, USA, June 28 - July 2, 2011*, pages 465–472. Omnipress, 2011.
- M. Gheshlaghi Azar, R. Munos, and H. J. Kappen. Minimax pac bounds on the sample complexity of reinforcement learning with a generative model. *Machine learning*, 91(3):325–349, 2013.
- M. H. Hado van Hasselt, Diana Borsa. Reinforcement learning lecture series 2021.
<https://deepmind.com/learning-resources/reinforcement-learning-series-2021>, 2015.
- A. Jonsson, E. Kaufmann, P. Ménard, O. Darwiche Domingues, E. Leurent, and M. Valko. Planning in markov decision processes with gap-dependent sample complexity. *Advances in Neural Information Processing Systems*, 33:1253–1263, 2020.
- C. Nota and P. S. Thomas. Is the policy gradient a gradient? In *Proceedings of the 19th International Conference on Autonomous Agents and MultiAgent Systems*, pages 939–947, 2020.
- D. Silver. Introduction to reinforcement learning.
<https://deepmind.com/learning-resources/-introduction-reinforcement-learning-david-silver>, 2015.



References II

- R. S. Sutton. Integrated architectures for learning, planning, and reacting based on approximating dynamic programming. In *Machine learning proceedings 1990*, pages 216–224. Elsevier, 1990.
- R. S. Sutton. Planning by incremental dynamic programming. In *Machine learning proceedings 1991*, pages 353–357. Elsevier, 1991.
- R. S. Sutton and A. G. Barto. *Reinforcement learning: An introduction*. MIT press, 2018.
- Y. Wan, M. Zaheer, A. White, M. White, and R. S. Sutton. Planning with expectation models. In *Proceedings of the 28th International Joint Conference on Artificial Intelligence*, pages 3649–3655, 2019.

